



# Data Engineering: Introduction to Stream Processing

## ■ Traditional Data Processing - Challenges

- Introduces too much “decision latency”
- Responses are delivered “after the fact”
- Maximum value of the identified situation is lost
- Decisions are made on old and stale data
- “Data at Rest”



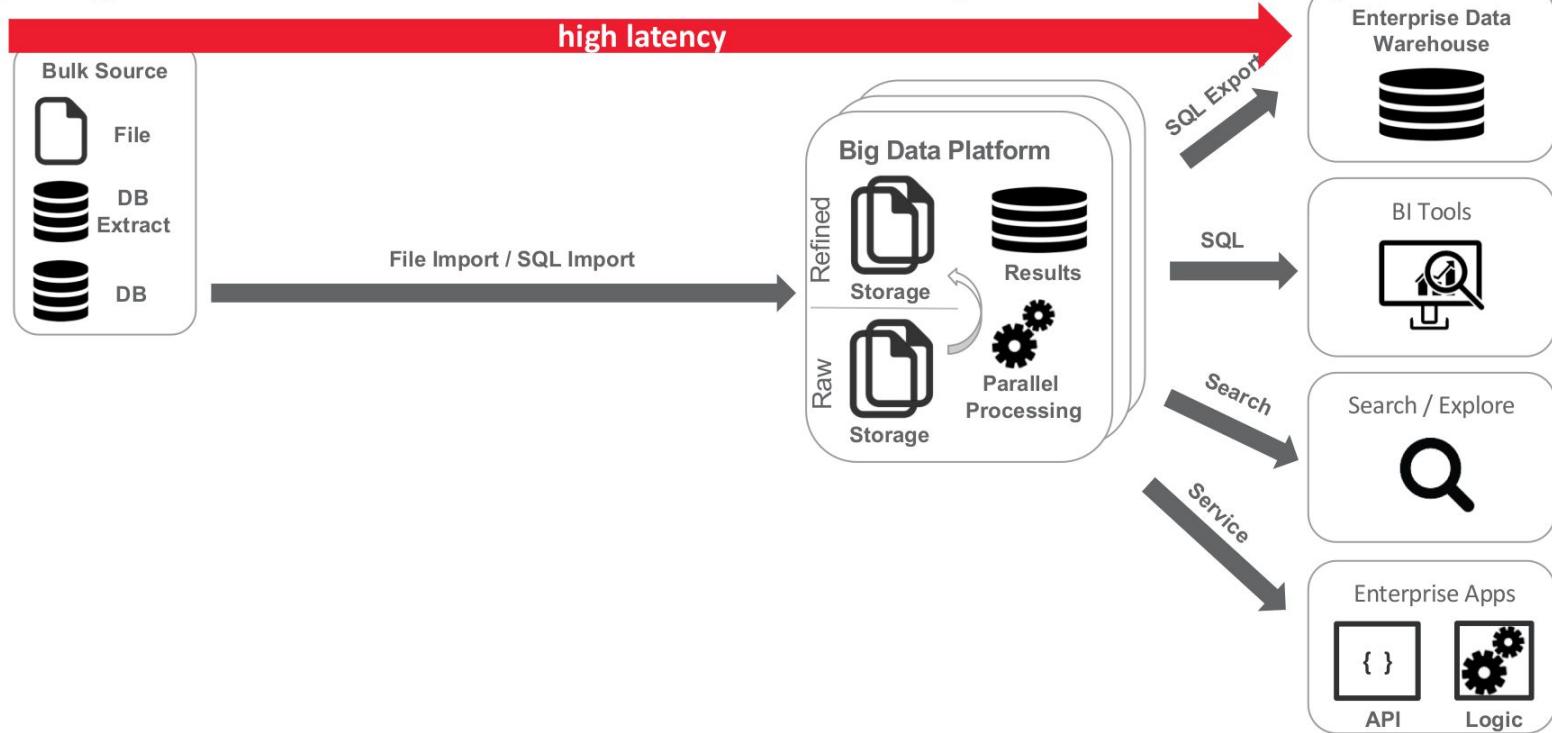
Introduction to Streaming Analytics



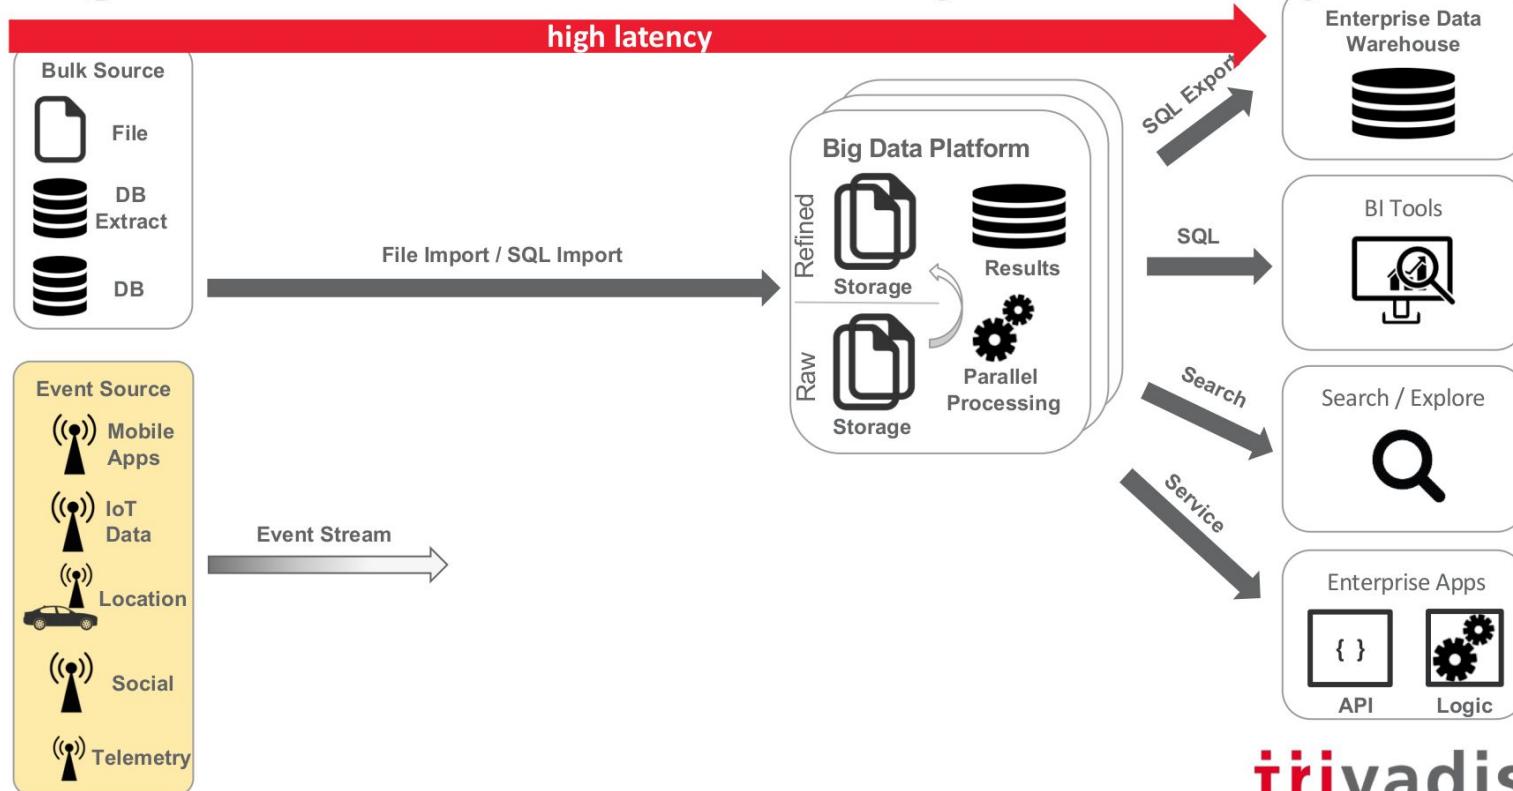
Data Engineering  
Introduction to Stream Processing

**trivadis**  
makes IT easier.

# ■ Big Data solves Volume and Variety – not Velocity



# ■ Big Data solves Volume and Variety – not Velocity



**trivadis**

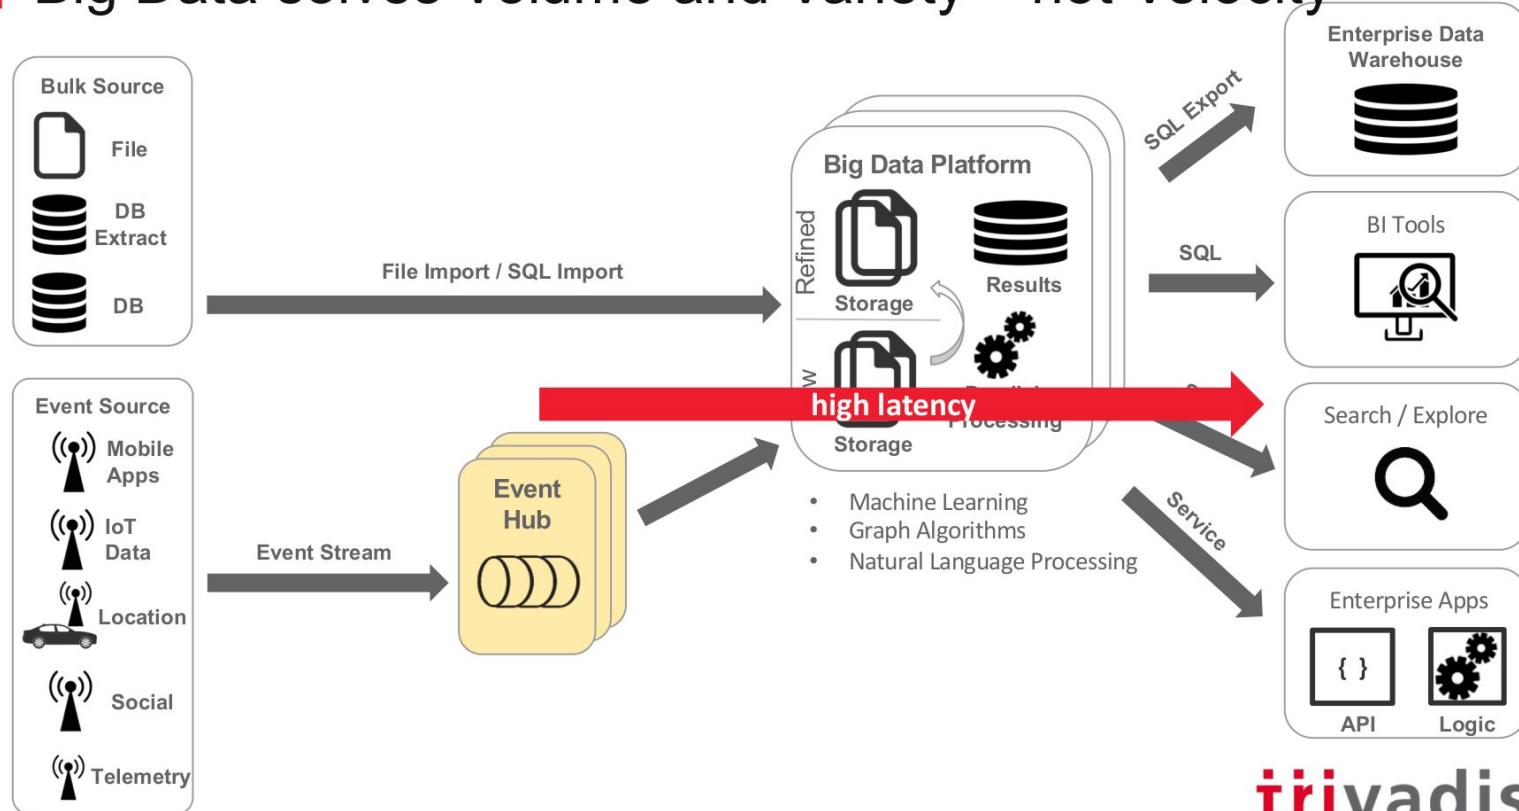


## ■ Real Time Analytics Use Cases

- Algorithmic Trading
- Online Fraud Detection
- Geo Fencing
- Proximity/Location Tracking
- Intrusion detection systems
- Traffic Management
- Recommendations
- Churn detection
- Internet of Things (IoT) / Intelligence Sensors
- Social Media/Data Analytics
- Gaming Data Feed
- ...



# ■ Big Data solves Volume and Variety – not Velocity

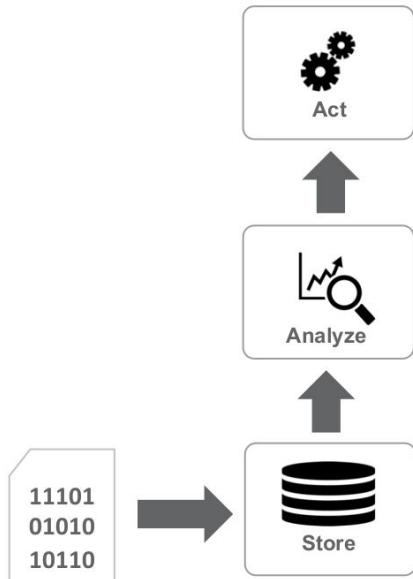


trivadis

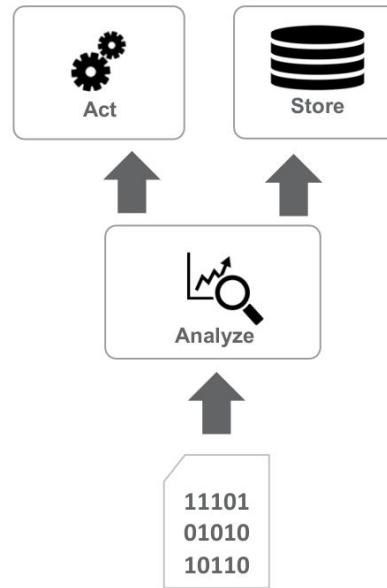


# ■ "Data at Rest" vs. "Data in Motion"

Data at Rest



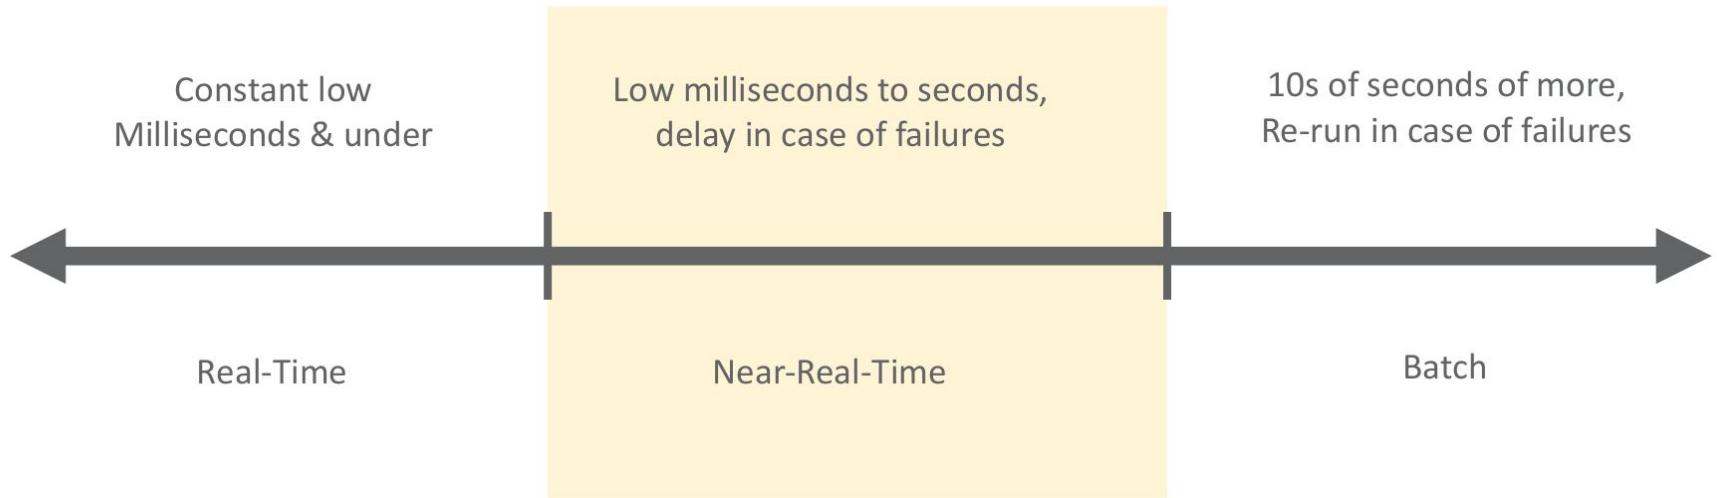
Data in Motion



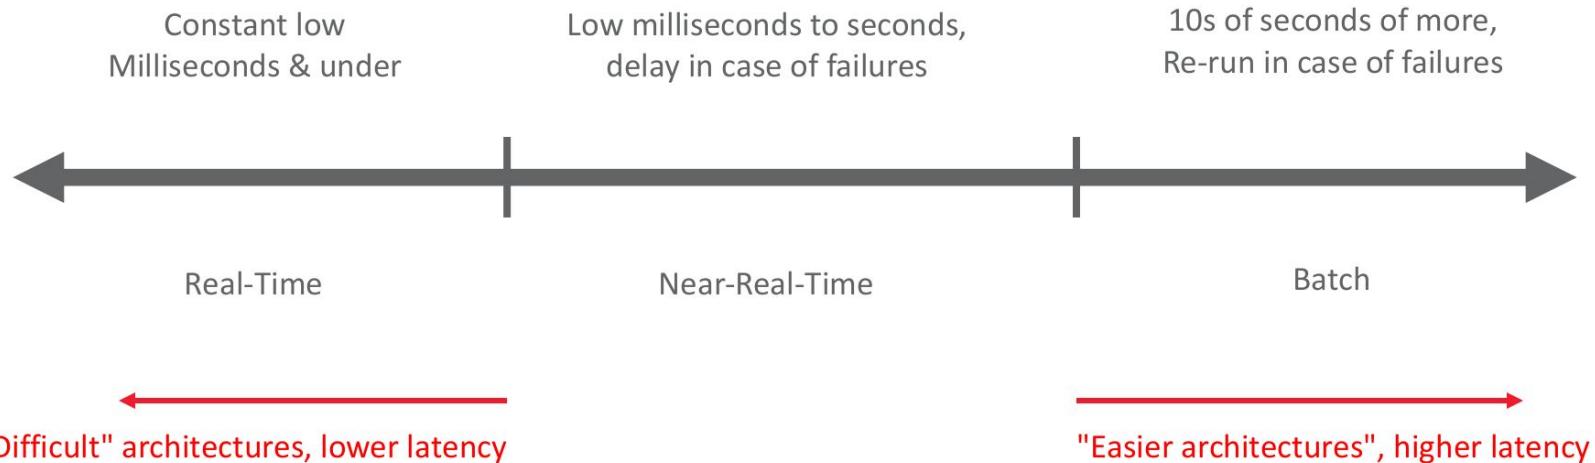
trivadis



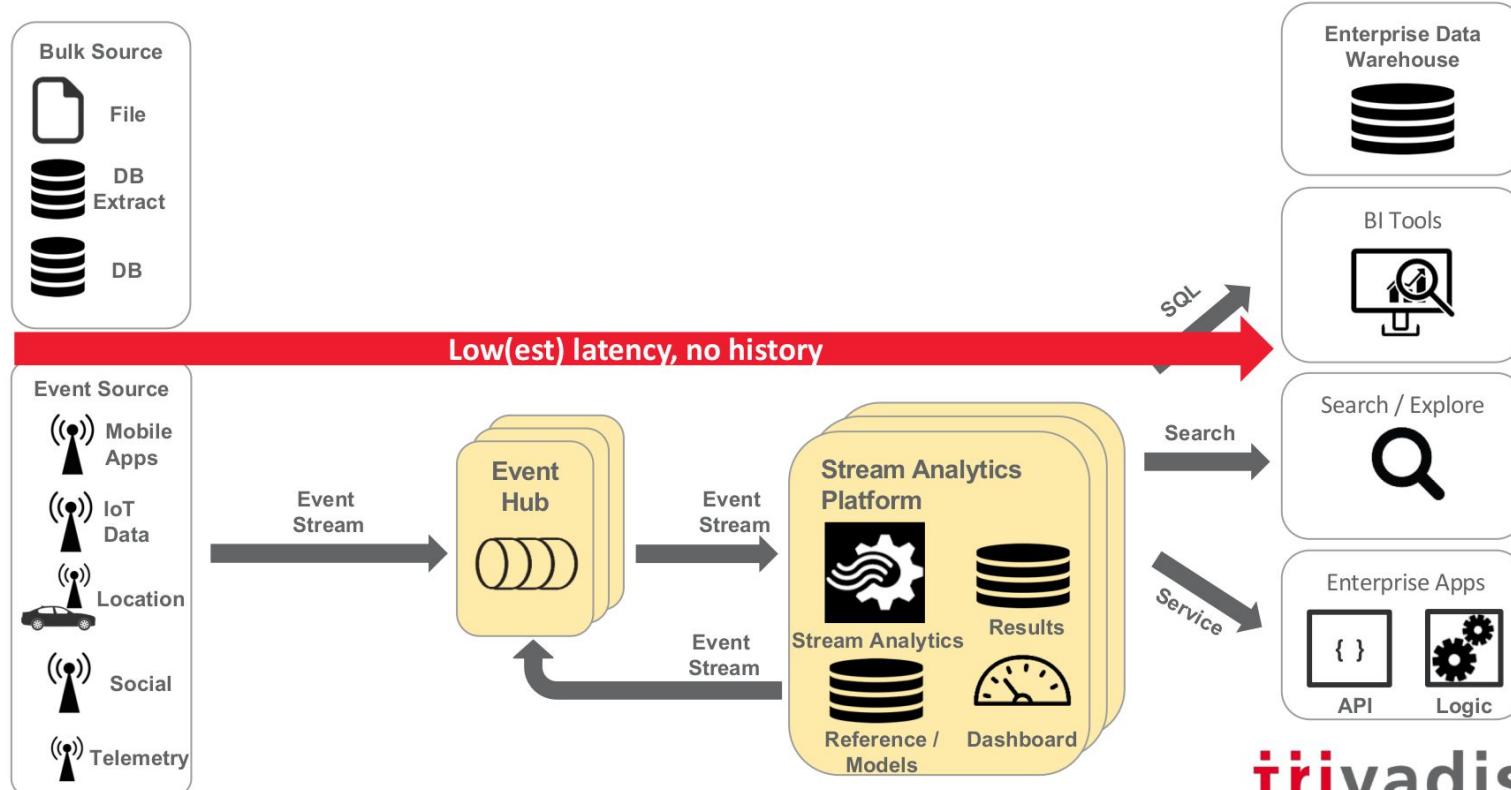
# ■ When to Stream / When not?



## ■ "No free lunch"



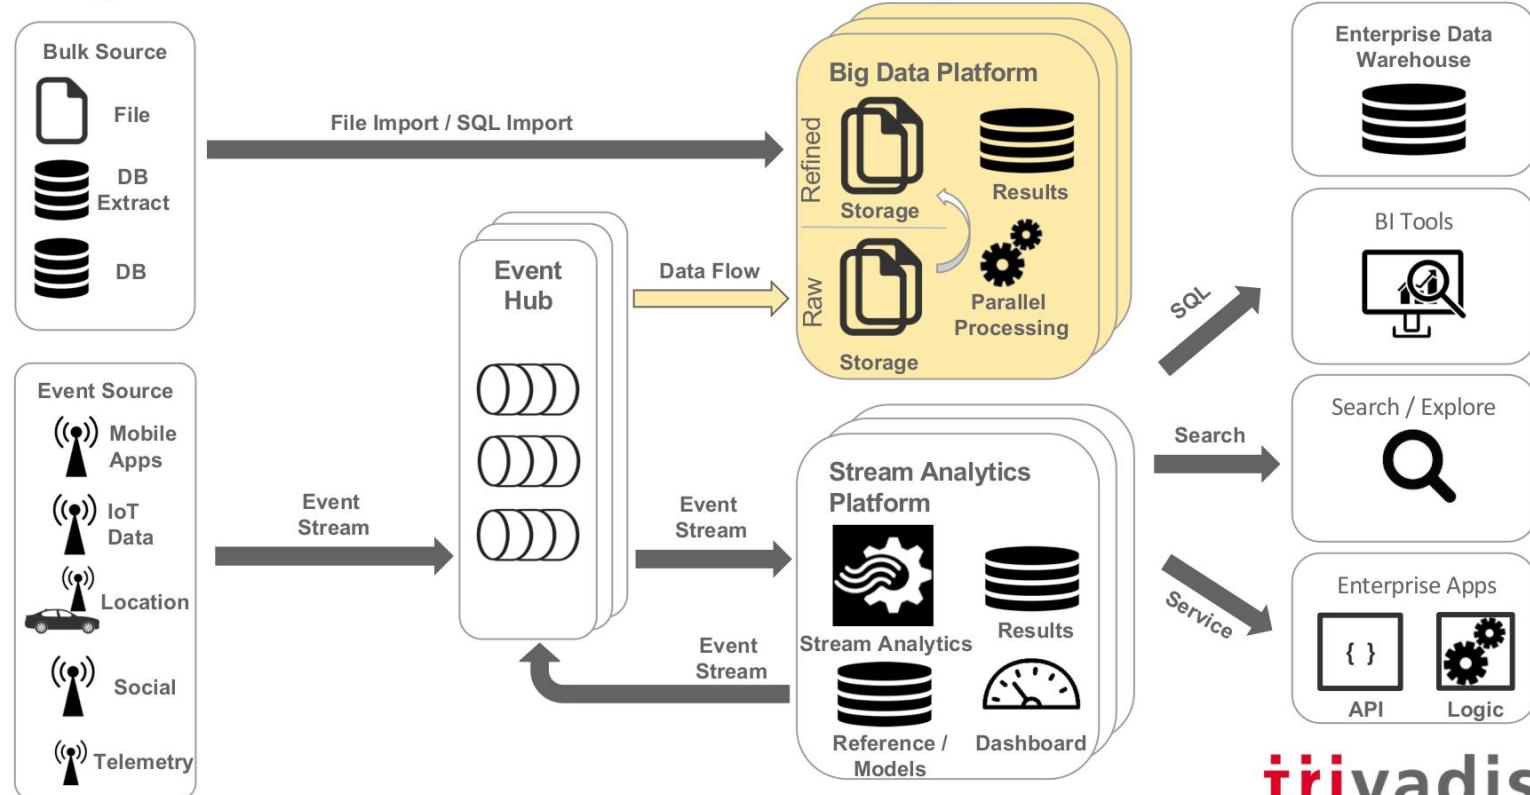
# ■ Stream Processing Architecture solves Velocity



trivadis



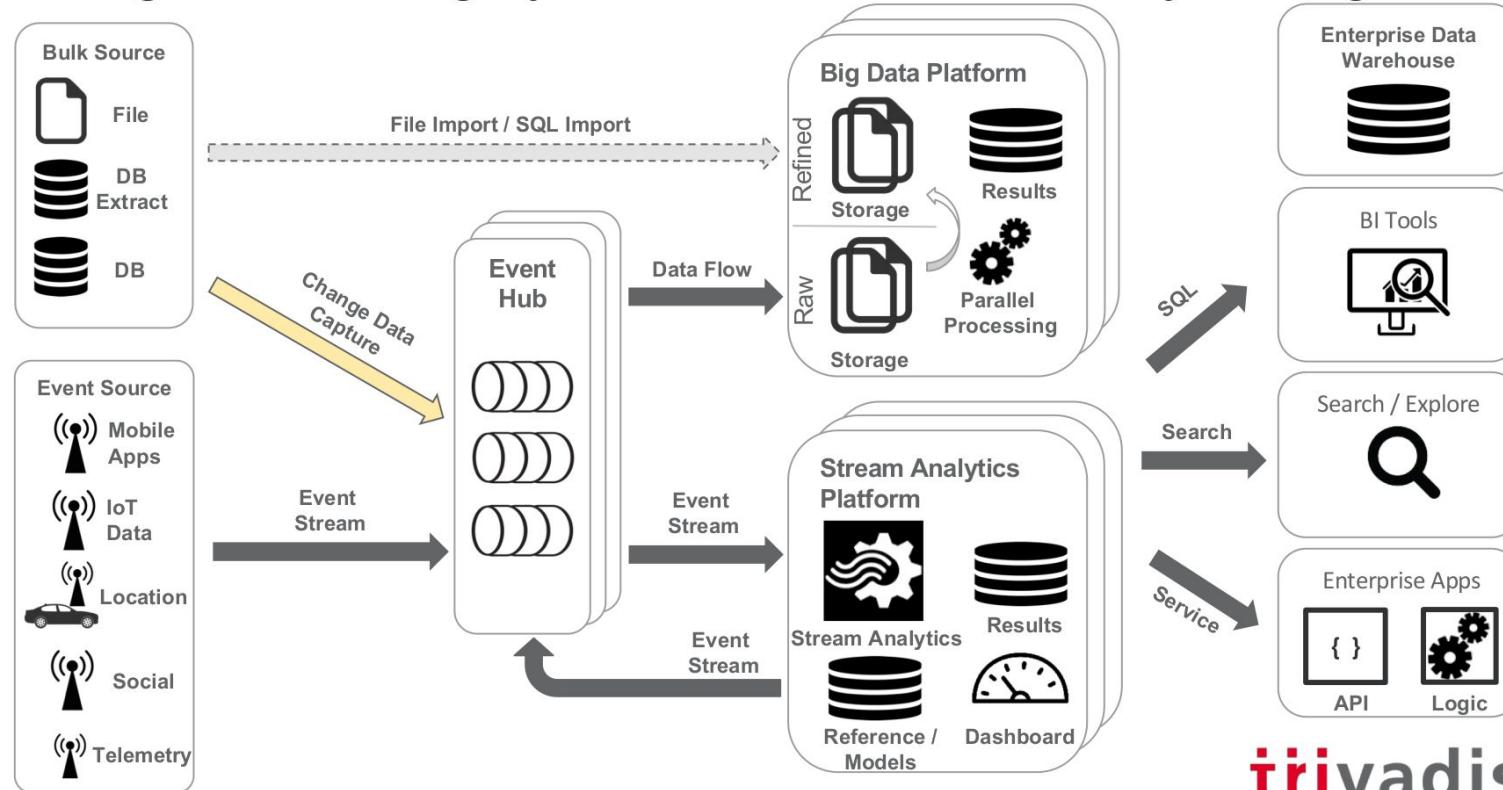
# ■ Big Data for all historical data analysis



trivadis



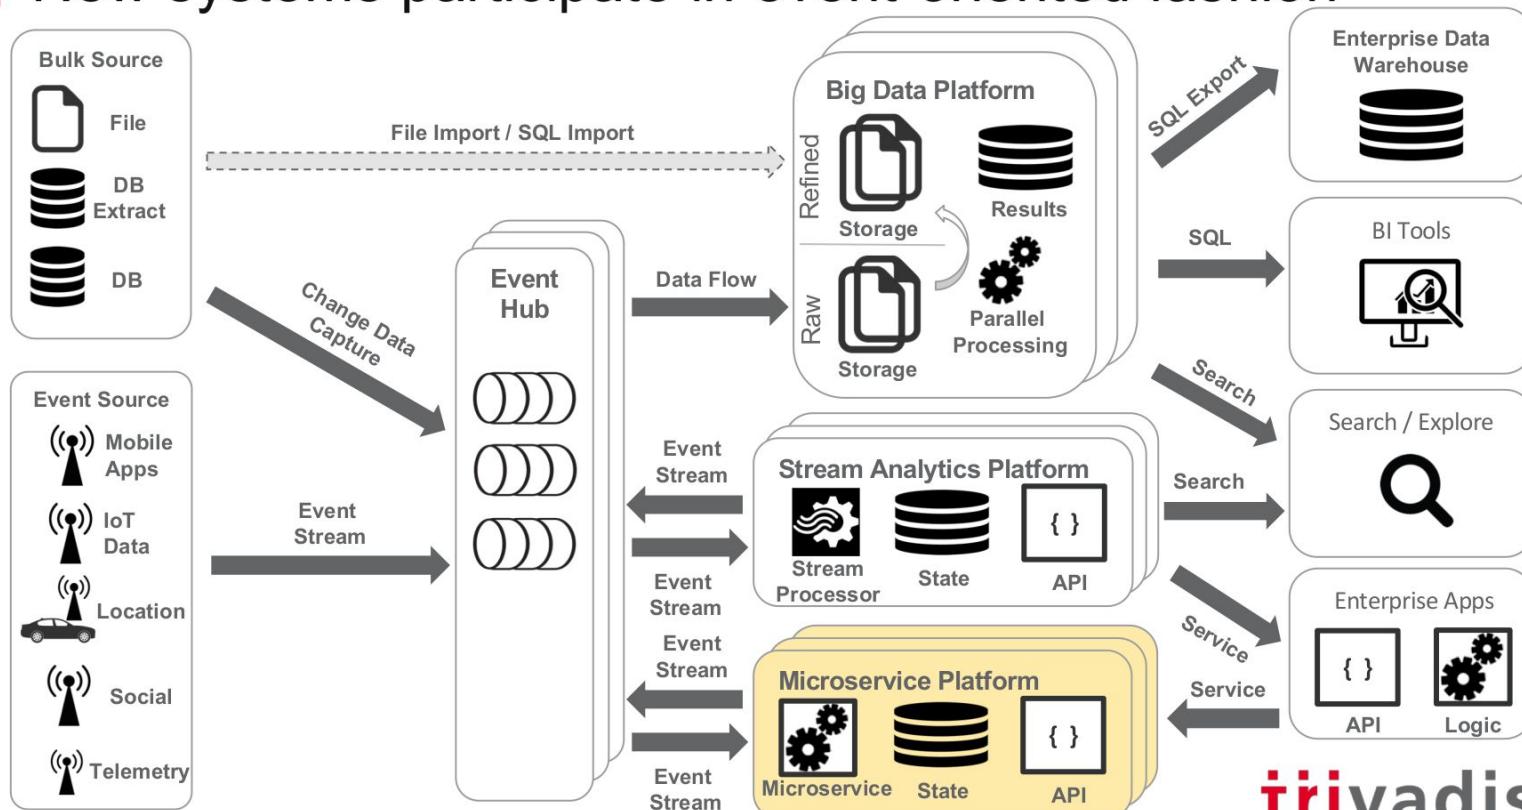
# ■ Integrate existing systems with lower latency through CDC



**trivadis**



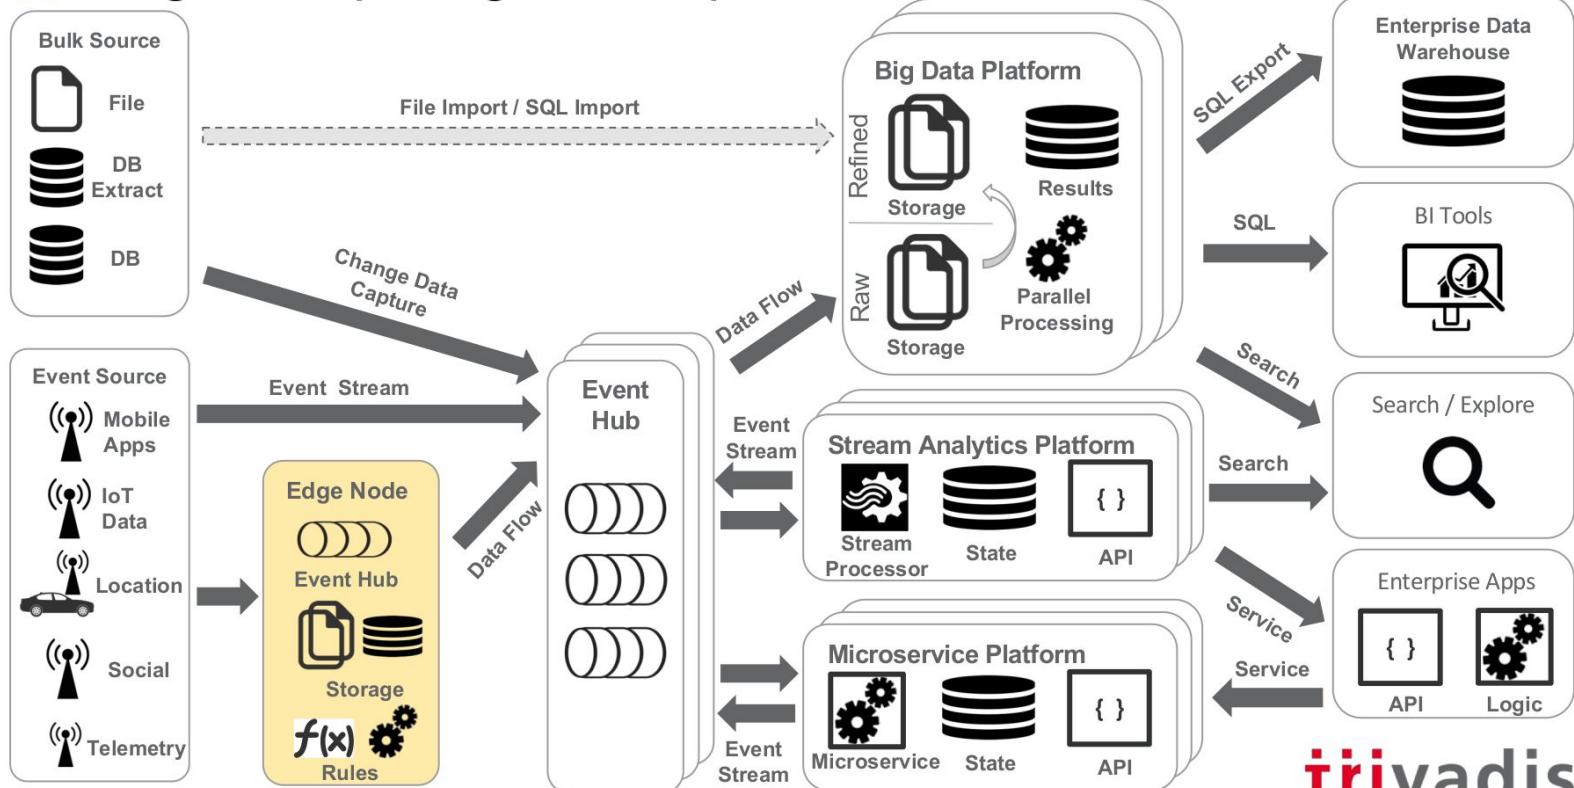
# ■ New systems participate in event-oriented fashion



trivadis



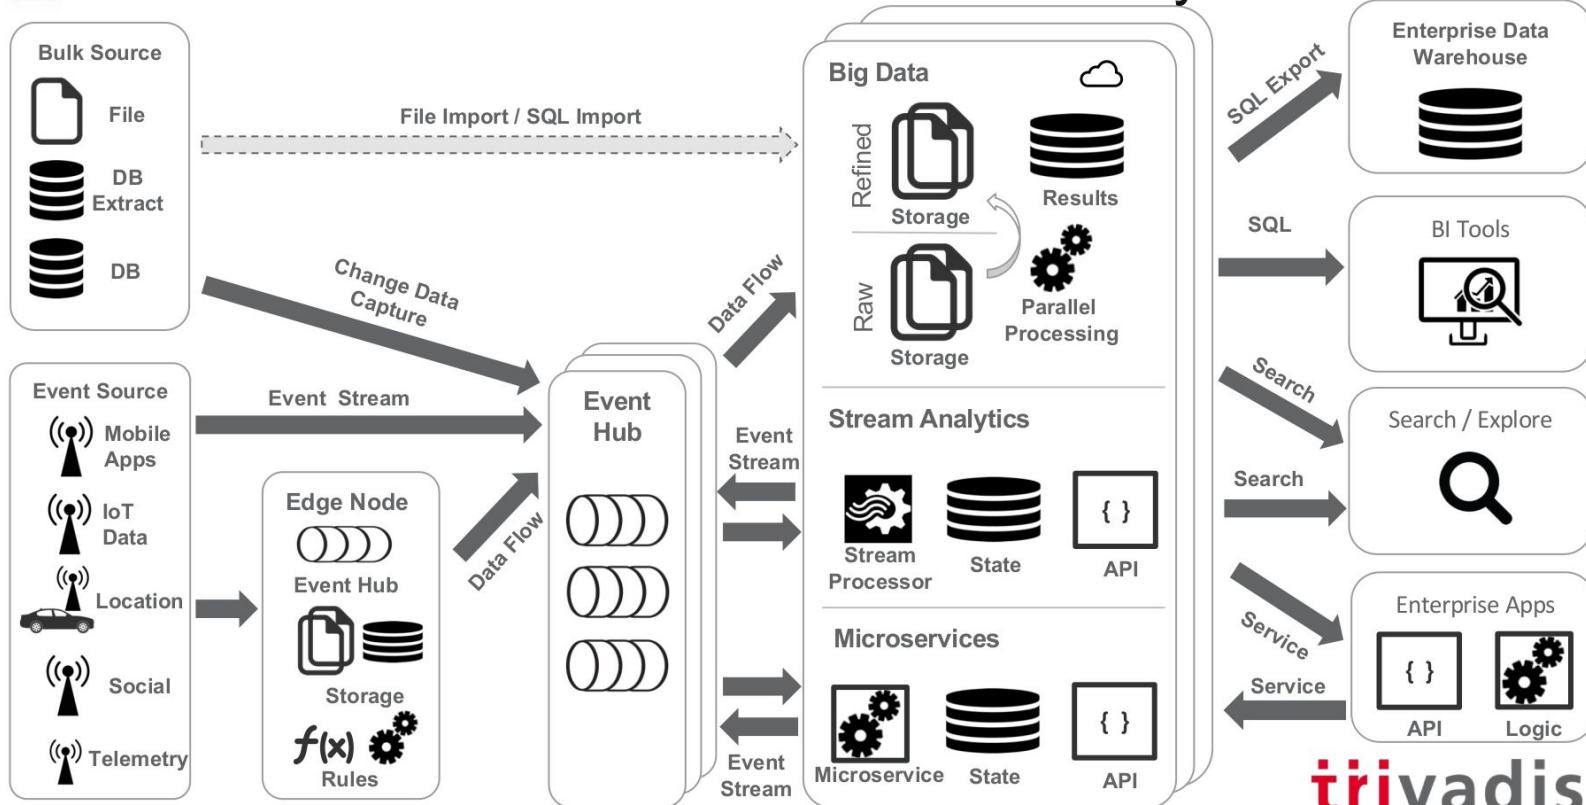
# ■ Edge computing allows processing close to data sources



trivadis



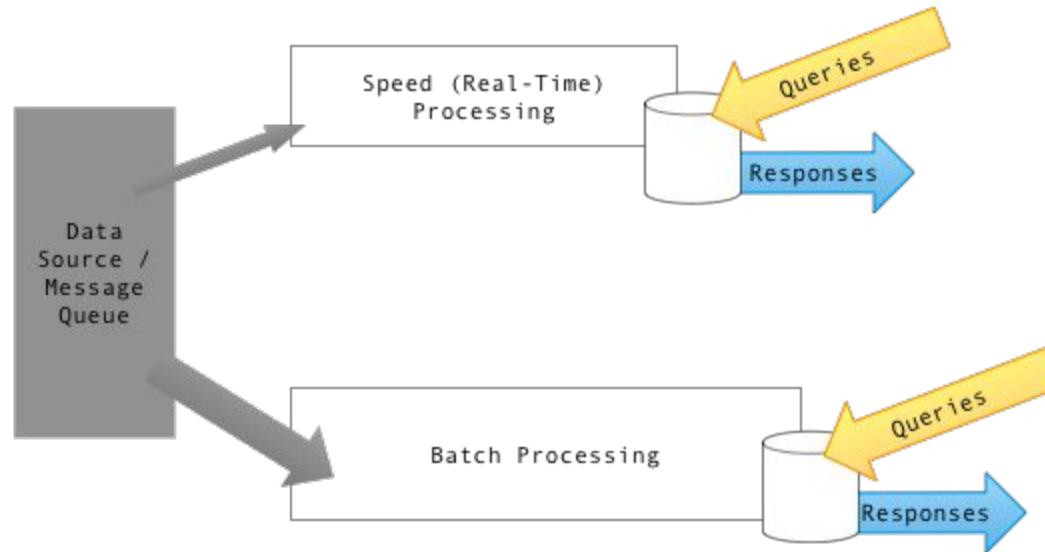
# ■ Unified Architecture for Modern Data Analytics Solutions



trivadis



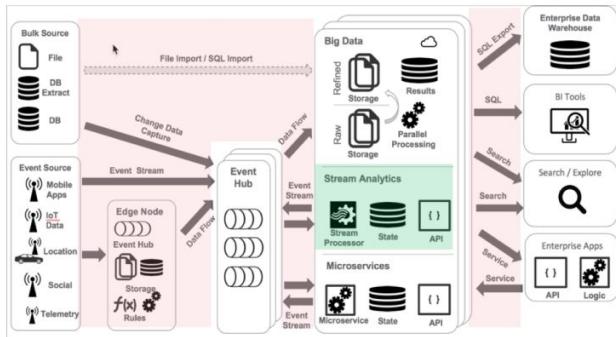
# Lambda Architecture



# ■ Two Types of Stream Processing (by Gartner)

## Stream Data Integration

- focuses on the ingestion and processing of data sources targeting real-time extract-transform-load (ETL) and data integration use cases
- filter and enrich the data



## Stream Analytics

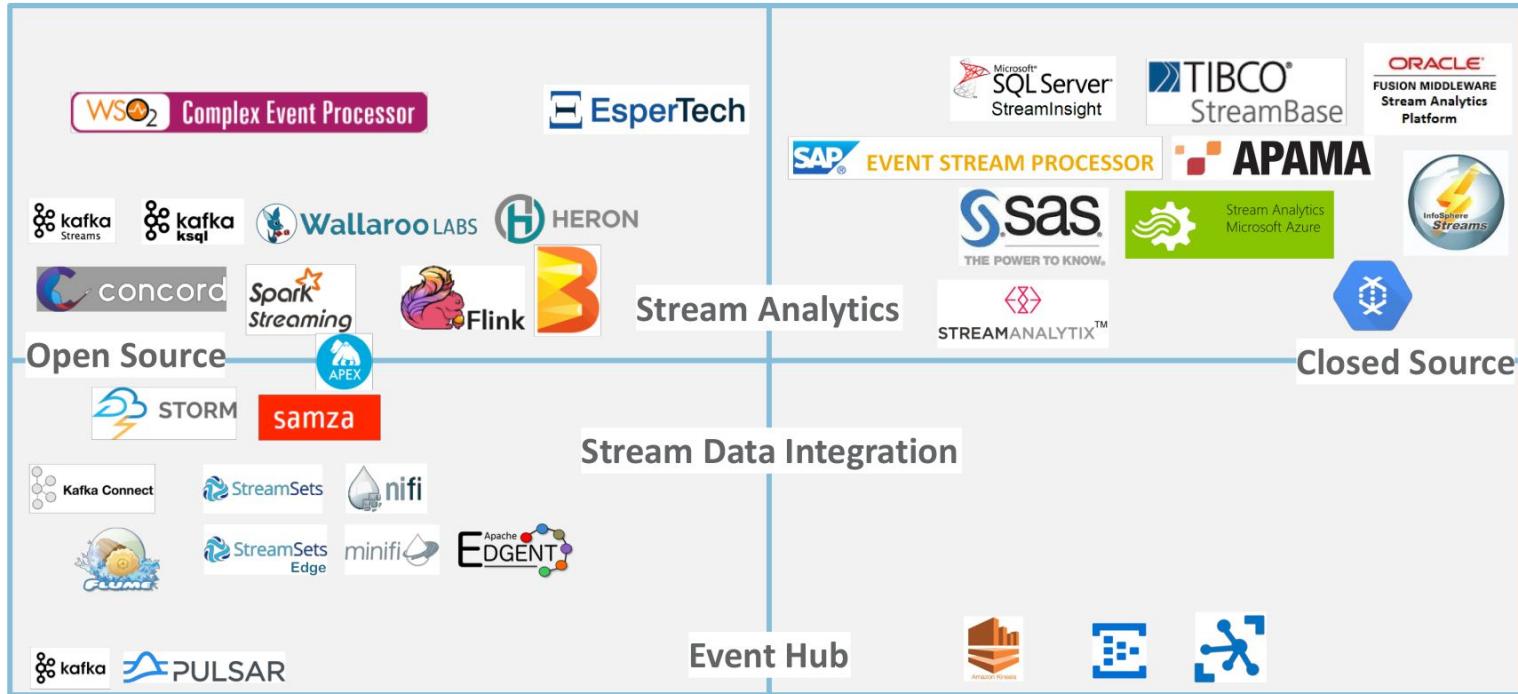
- targets analytics use cases
- calculating aggregates and detecting patterns to generate higher-level, more relevant summary information (complex events)
- Complex events may signify threats or opportunities that require a response from the business

Gartner: Market Guide for Event Stream Processing, Nick Heudecker, W. Roy Schulte

**trivadis**



# Stream Processing & Analytics Ecosystem



Source: adapted from Tibco

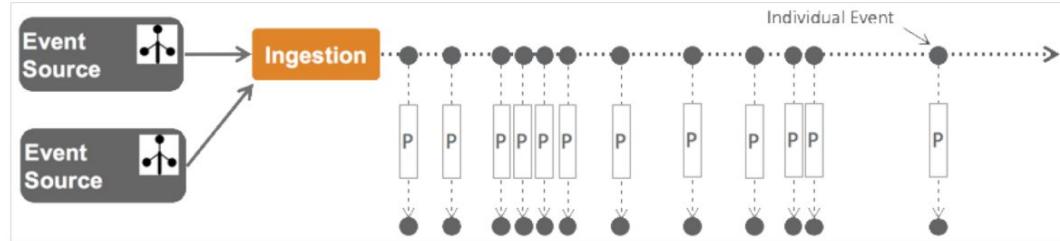
trivadis



## ■ Event-at-a-time vs. Micro Batch

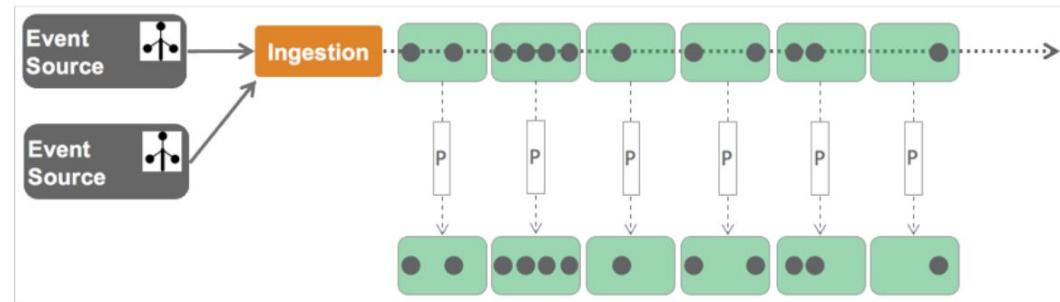
### Event-at-a-time Processing

- Events processed as they arrive
- low-latency
- fault tolerance expensive



### Micro-Batch Processing

- Splits incoming stream in small batches
- Fault tolerance easier
- Better throughput



# ■ Delivery Guarantees



## At most once (fire-and-forget)

message is sent, but the sender **doesn't care if it's received or lost**. it is *the easiest and most performant behavior to support*.



## At least once

retransmission of a message will occur until an acknowledgment is received. Since a delayed acknowledgment from a receiver could be in flight when the sender retransmits, the message may be received **one or more times**.



## Exactly once

ensures that a message is received **once and only once**, and is never lost and never repeated. The system must implement whatever mechanisms are required to ensure that a message is received and processed just once



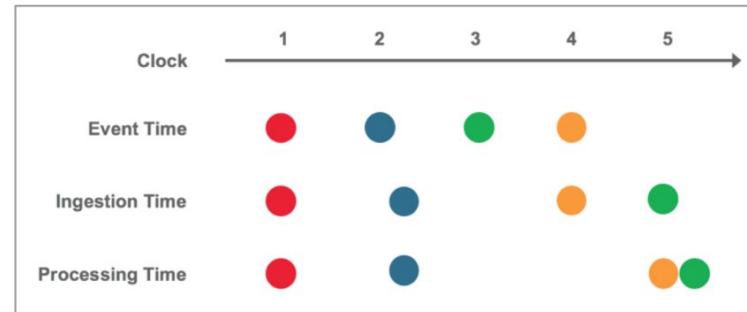
# ■ Event Time vs. Ingestion / Processing Time

## Event time

the time at which events actually *occurred*

## Ingestion time / Processing Time

the time at which events are *ingested into / processed by* the system



Not all use cases care about event times but many do

## Examples

- characterizing user behavior over time
- most billing applications
- anomaly detection



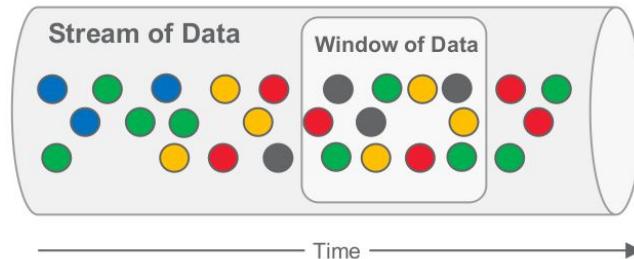
# ■ Windowing

Computations over events done using  
**windows of data**

Due to size and never-ending nature of it,  
it's not feasible to keep entire stream of  
data in memory

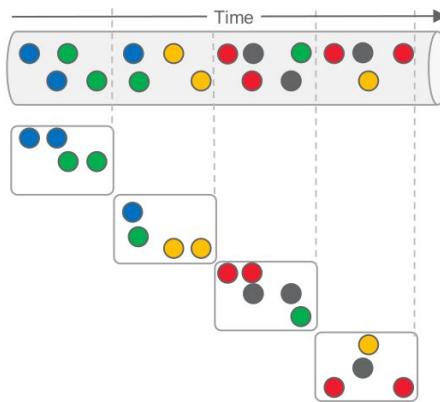
A window of data represents a certain  
amount of data where we can perform  
computations on

Windows give the power to keep a  
working memory and look back at recent  
data efficiently



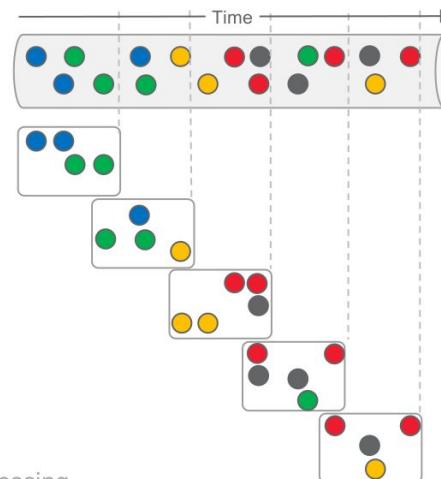
# ■ Windowing

**Fixed Window (aka Tumbling Window)** - eviction policy always based on the window being full and trigger policy based on either the count of items in the window or time

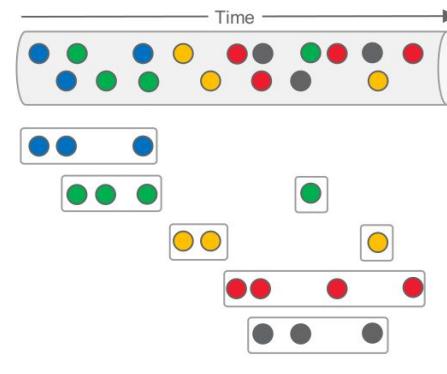


Introduction to Stream Processing

**Sliding Window (aka Hopping Window)** - uses eviction and trigger policies that are based on time: *window length* and *sliding interval length*



**Session Window** – composed of sequences of temporarily related events terminated by a gap of inactivity greater than some timeout



**trivadis**   
@gschmutz makes IT easier.

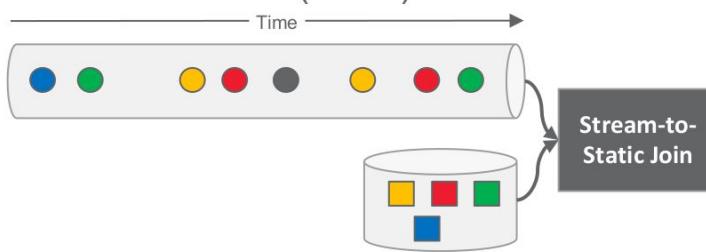


# ■ Joining – Stream-to-Static and Stream-to-Stream

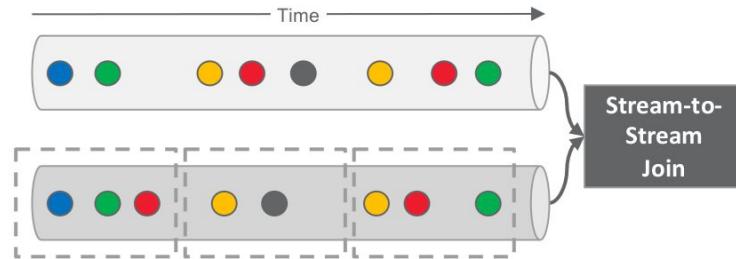
## Challenges of joining streams

1. Data streams need to be aligned as they come because they have different timestamps
2. since streams are never-ending, the joins must be limited; otherwise join will never end
3. join needs to produce results continuously as there is no end to the data

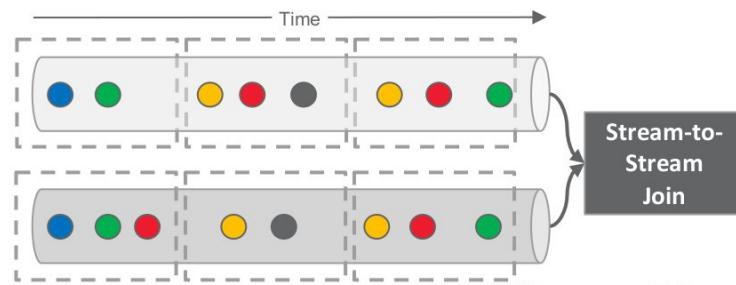
## Stream-to-Static (Table) Join



## Stream-to-Stream Join (one window join)



## Stream-to-Stream Join (two window join)



**trivadis**



# ■ State Management

Necessary if stream processing use case  
is dependent on previously seen data or  
external data

**Windowing, Joining and Pattern Detection** use State Management behind the scenes

State Management services can be made available for custom state handling logic

State needs to be managed as close to the stream processor as possible

## Options for State Management

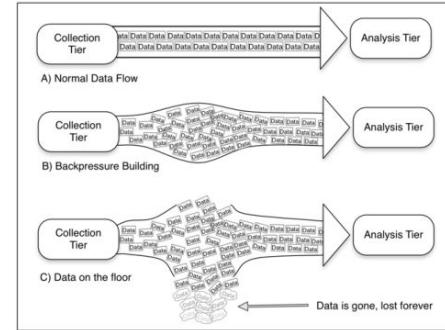


How does it handle failures? If a machine crashes and the/some state is lost?



## ■ 7) Back Pressure

- Backpressure refers to the situation where a system is receiving data at a higher rate than it can process
- Backpressure, if not dealt with correctly, can lead to exhaustion of resources, or even, in the worst case, data loss



- A slow consumer should slow down the producer

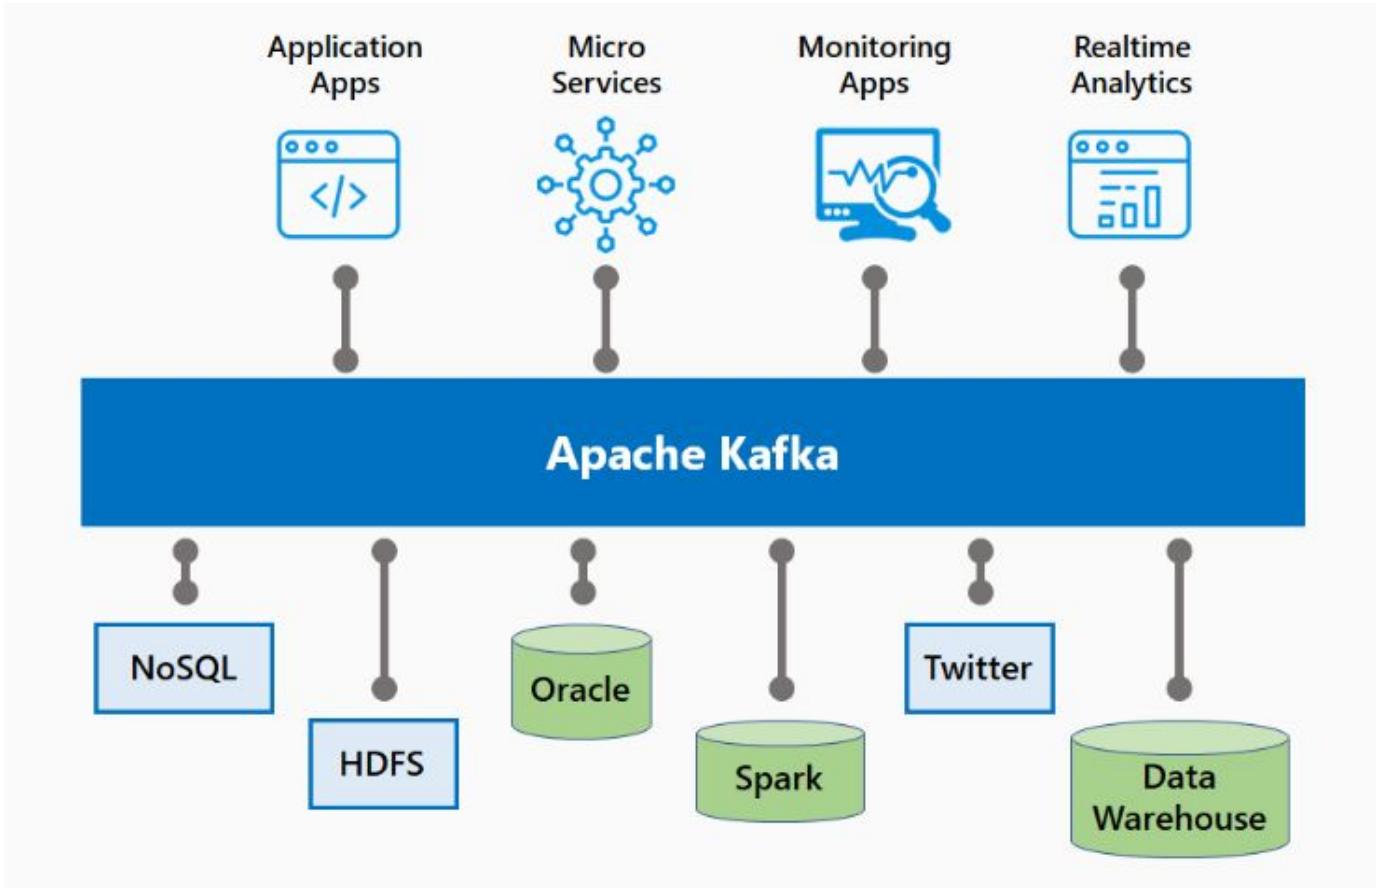


Introduction to Streaming Analytics

**trivadis**  
makes IT easier. ■ ■ ■







*“Kafka is a distributed, partitioned, replicated commit log service. It provides the functionality of a messaging system, but with a unique design”*

## History

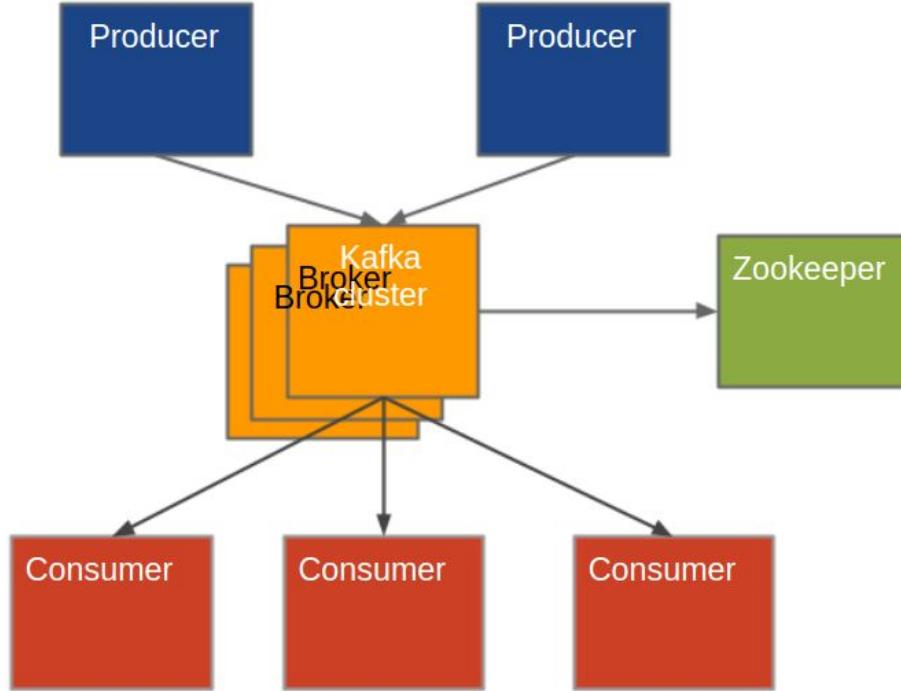
- Apache Kafka was originally developed by LinkedIn, open sourced in early 2011
- Written in Scala



# Apache Kafka

- **Fast:** A single Kafka broker can handle hundreds of megabytes of reads and writes per second from thousands of clients.
- **Scalable:** Kafka is designed to allow a single cluster to serve as the central data backbone for a large organization. It can be elastically and transparently expanded without downtime. Data streams are partitioned and spread over a cluster of machines to allow data streams larger than the capability of any single machine and to allow clusters of co-ordinated consumers
- **Durable:** Messages are persisted on disk and replicated within the cluster to prevent data loss. Each broker can handle terabytes of messages without performance impact.
- **Distributed by Design:** Kafka has a modern cluster-centric design that offers strong durability and fault-tolerance guarantees.

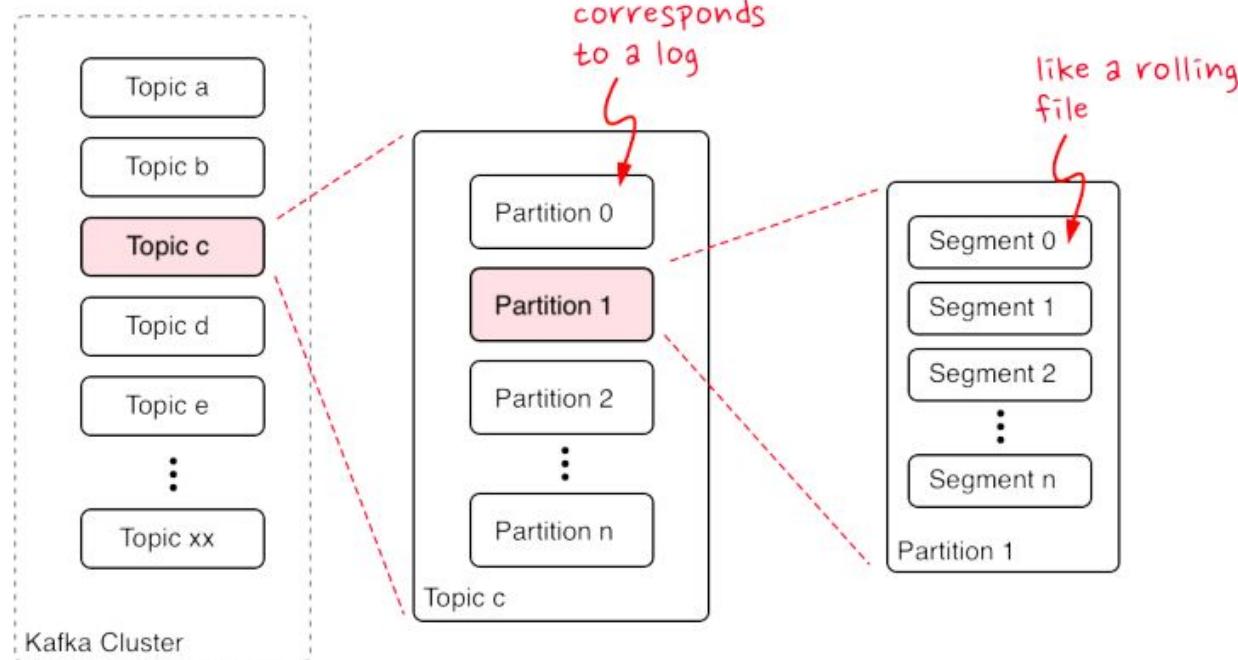




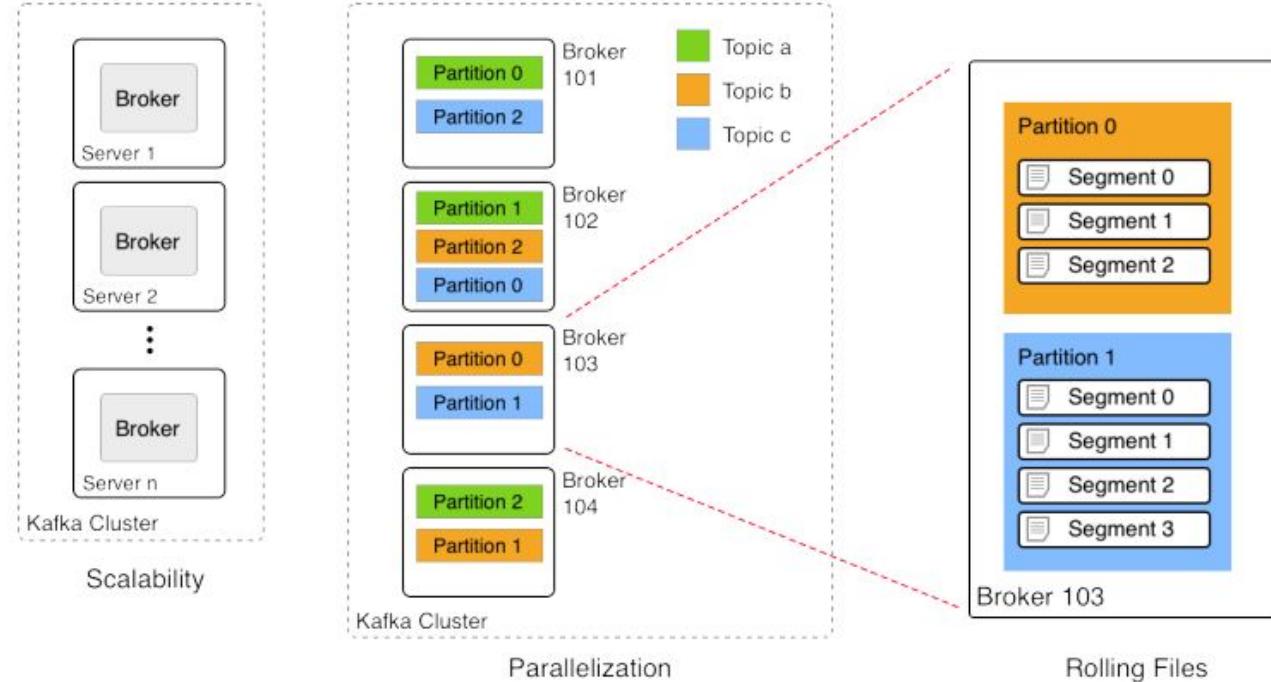
- *Producers* publish messages to a Kafka topic
- *Consumers* process the feed of published messages
- *Cluster* comprises one or more servers (broker)



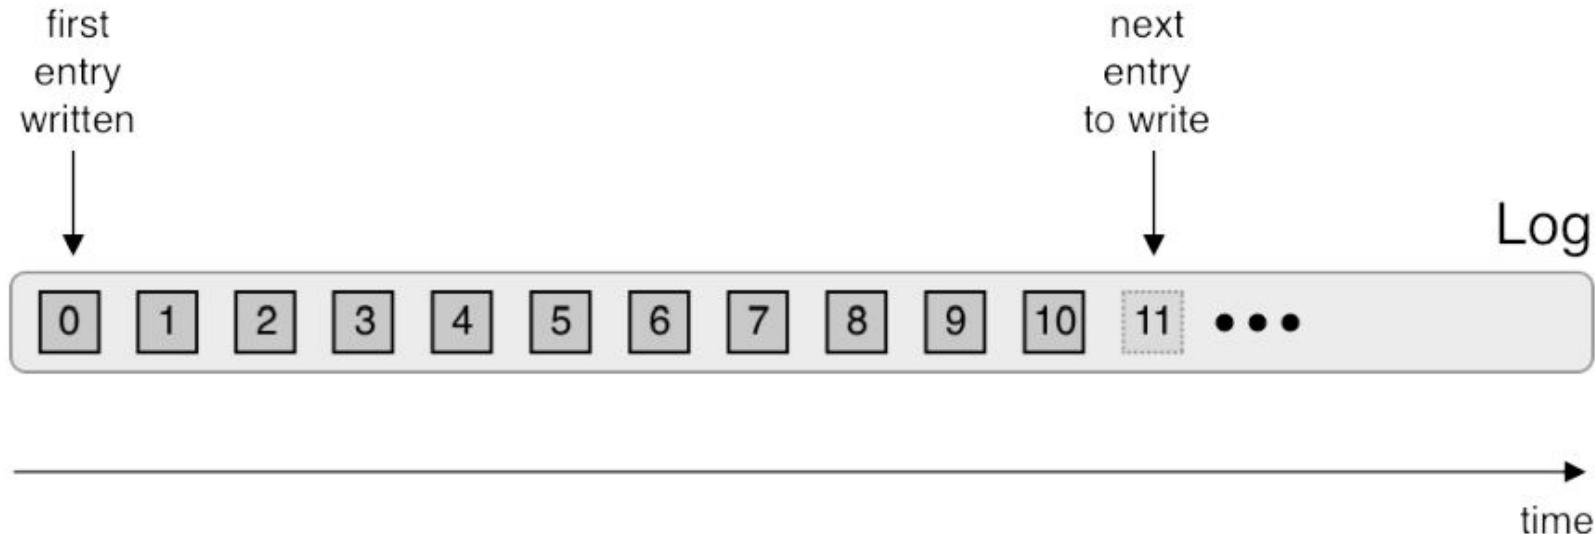
# Topics, Partitions, and Segments



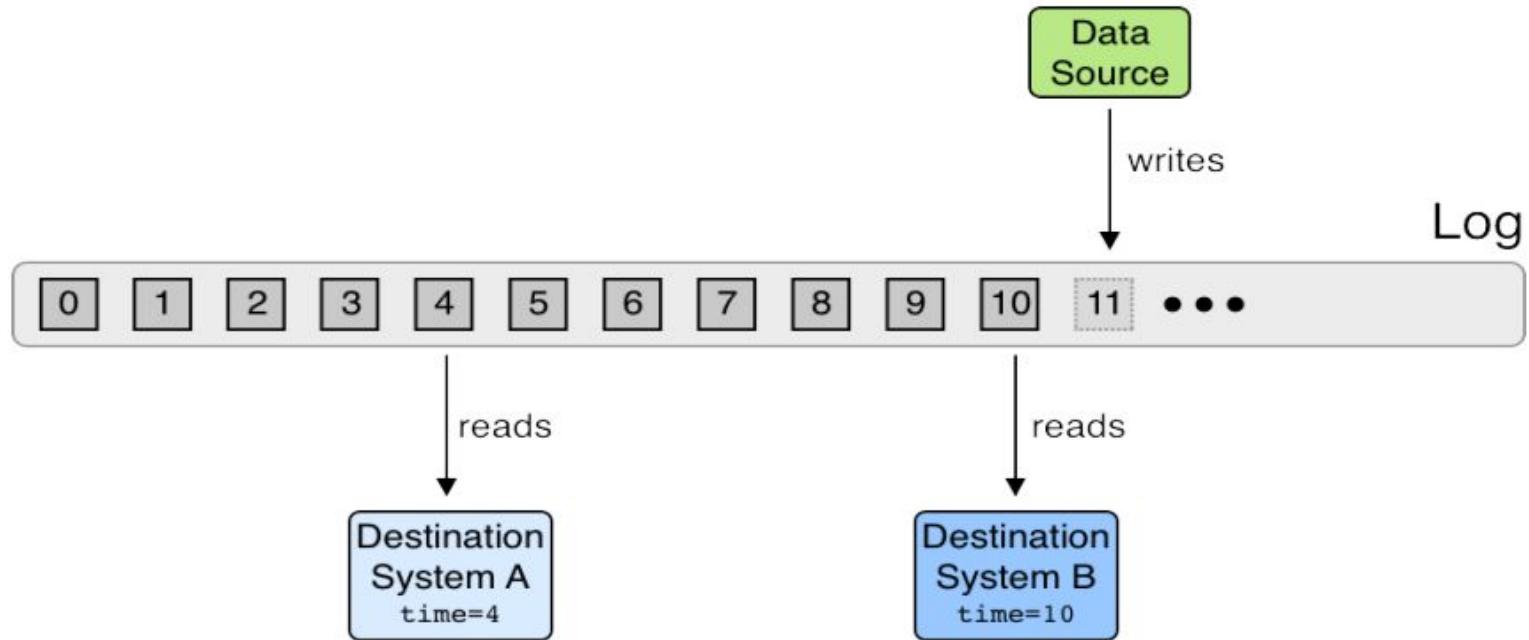
# Topics, Partitions, and Segments



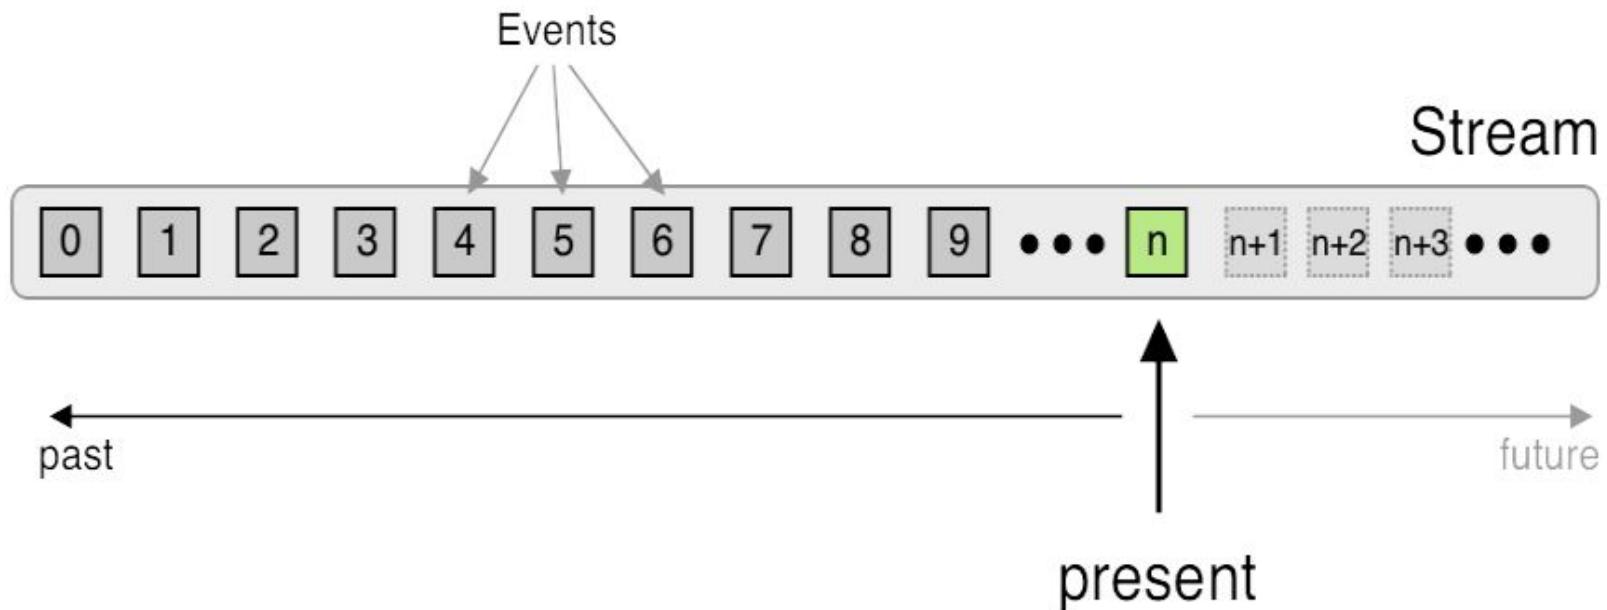
# The Log



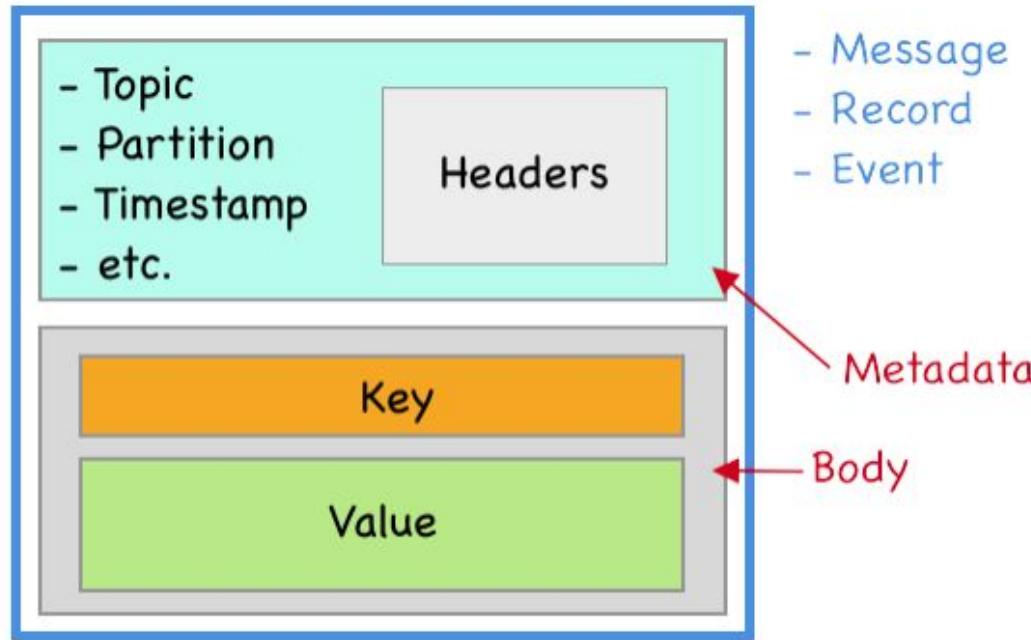
# Log Structured Data Flow



# The Stream



# Data Elements



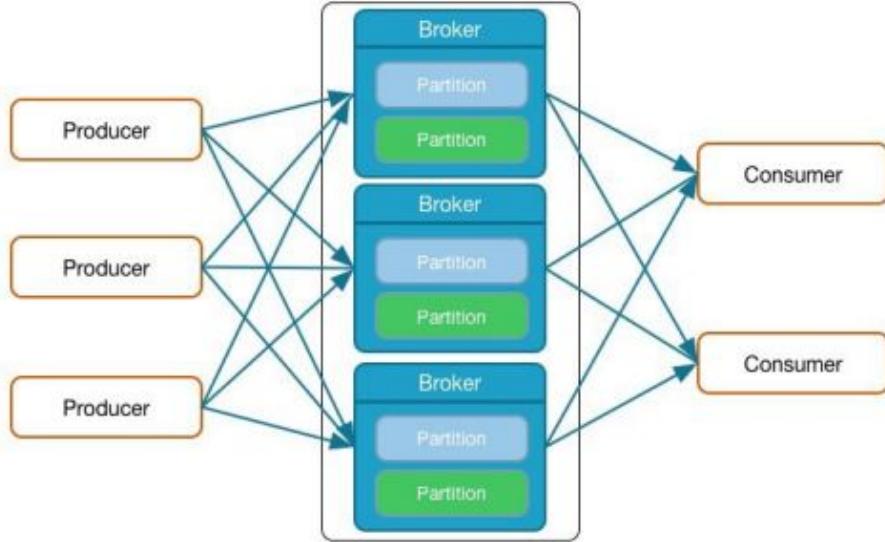
# Brokers Manage Partitions

- Messages of Topic spread across Partitions
- Partitions spread across Brokers
- Each Broker handles many Partitions
- Each Partition stored on Broker's disk
- Partition: 1..n log files
- Each message in Log identified by *Offset*
- Configurable Retention Policy

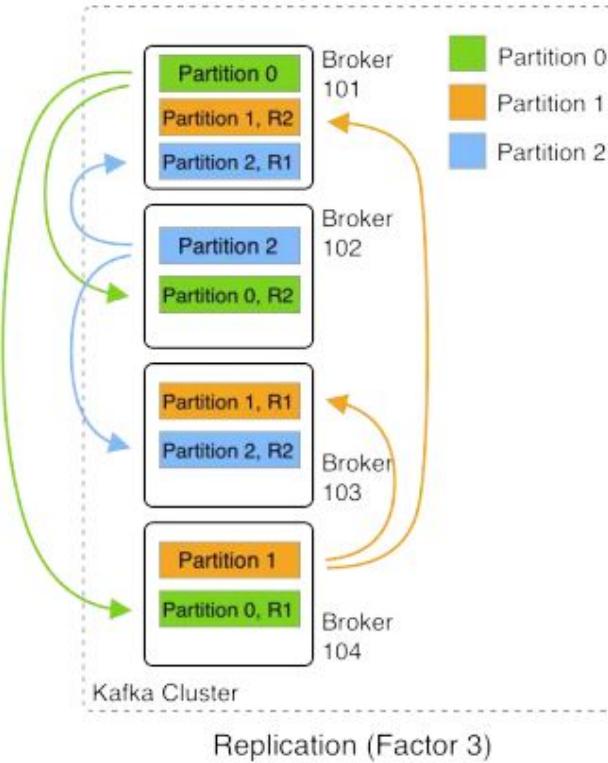


# Broker Basics

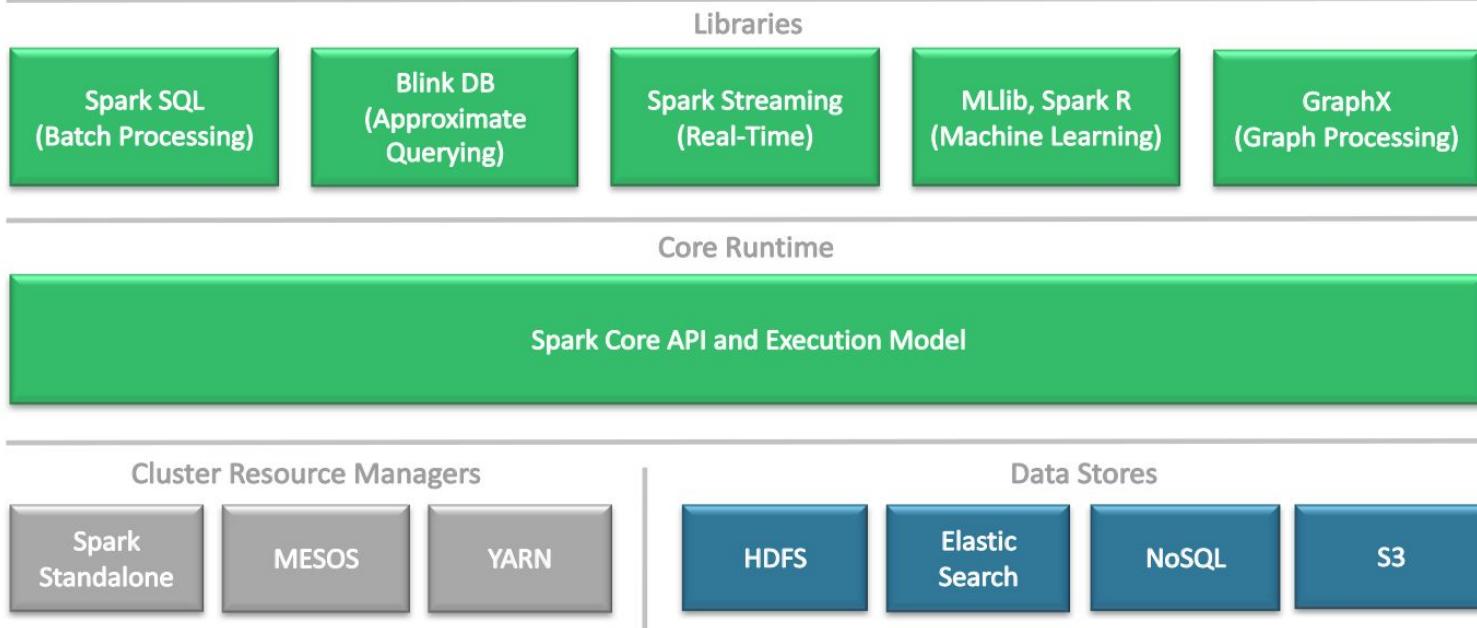
- Producer sends Messages to Brokers
- Brokers receive and store Messages
- A Kafka Cluster can have many Brokers
- Each Broker manages multiple Partitions



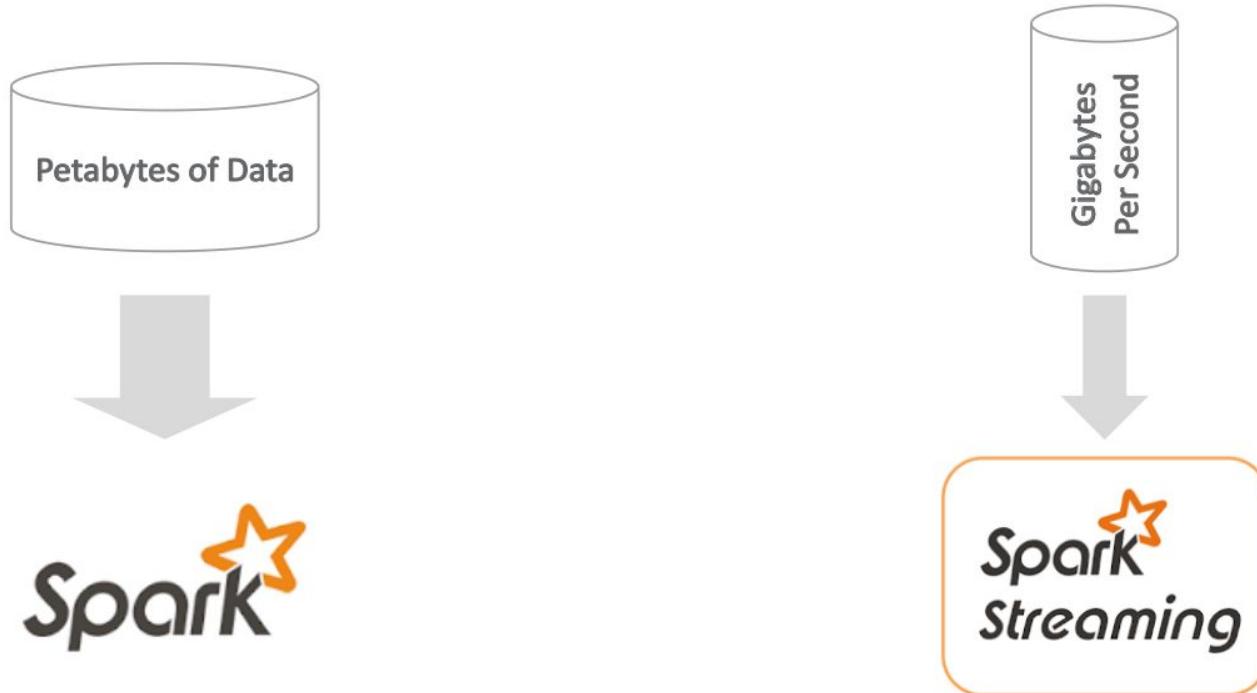
# Broker Replication



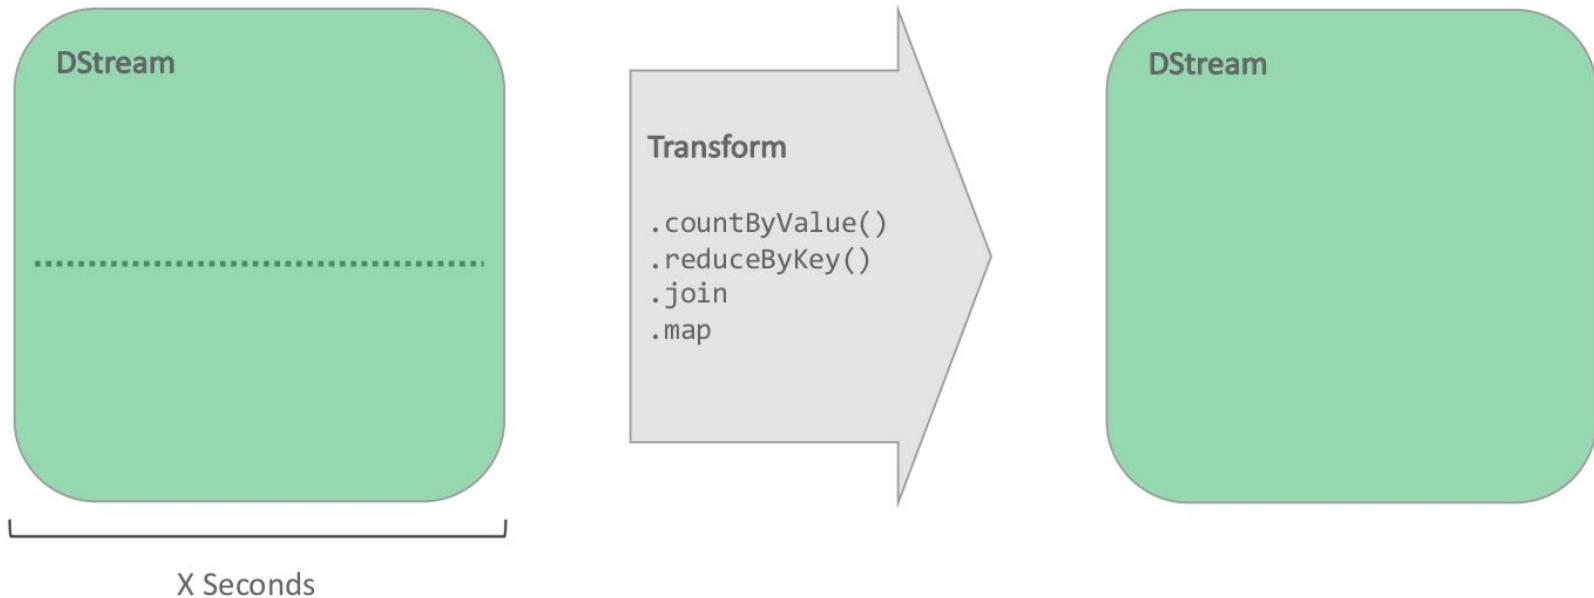
## ■ Apache Spark – Stack



## ■ Apache Spark - Batch vs. Real-Time Processing



# ■ Apache Spark Streaming



# Apache Spark Streaming



# Apache Spark Streaming

```
class TwitterStreamingHashTagsByInterval extends Serializable {

    def start(auth: Option[Authorization], ssc: StreamingContext, filters: Regex, keyspace: String, table: String): Unit = {

        val transform = (cruft: String) => filters.findAllIn(cruft).flatMap(_.stripPrefix("#"))

        val stream = TwitterUtils.createStream(ssc, auth, Nil, StorageLevel.MEMORY_ONLY_SER_2)

        /* Note that Cassandra is doing the sorting for you here. */
        stream.flatMap(_.getText.toLowerCase.split("""\s+"""))
            .map(transform)
            .countByValueAndWindow(Seconds(5), Seconds(5))
            .transform((rdd, time) => rdd.map { case (term, count) => (term, count, now(time)) })
            .saveToCassandra(keyspace, table, SomeColumns("hashtag", "mentions", "interval"))

        ssc.checkpoint("./checkpoint")
        ssc.start()
        ssc.awaitTermination()
    }

    private def now(time: Time): String =
        new DateTime(time.milliseconds, DateTimeZone.UTC).toString("yyyyMMddHH:mm:ss.SSS")
}
}
```

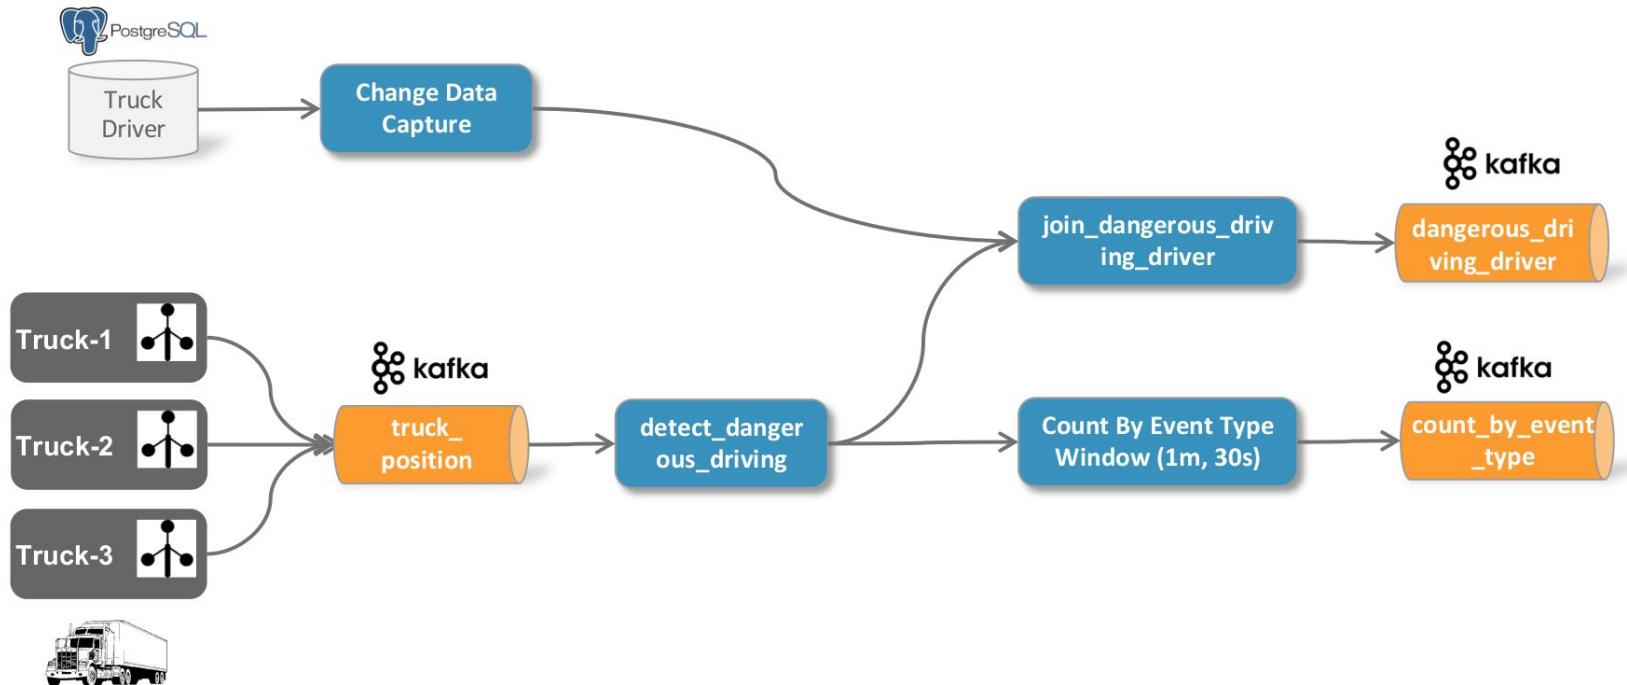


```
val orderDf = spark
    .readStream
    .format("kafka")
    .option("kafka.bootstrap.servers",
"broker-1:9092")
    .option("subscribe", "order")
    .load()

val orderFilteredDf =
    orderDf.where(
        "address.county = 'Switzerland'")
```

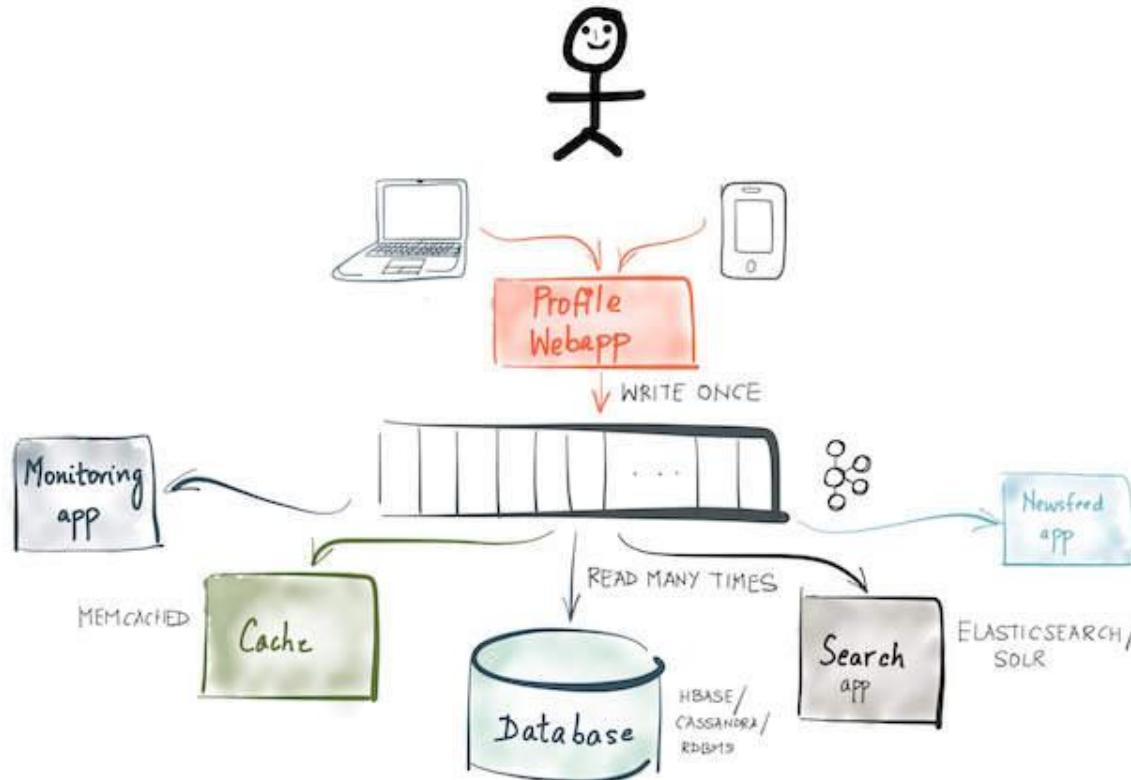


# Sample Use Case

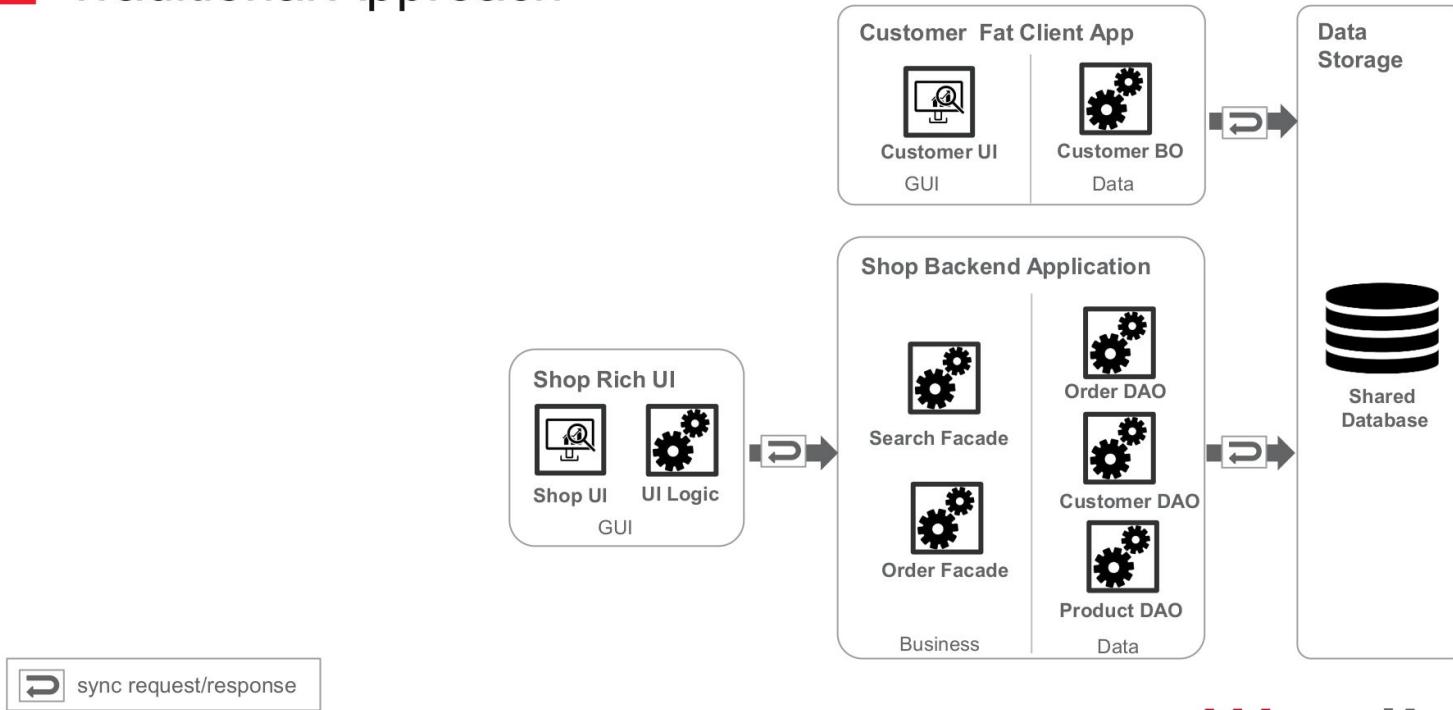


# Event Sourcing





# ■ Traditional Approach



Building event-driven Microservices with Kafka Ecosystem

**trivadis**  
makes IT easier.



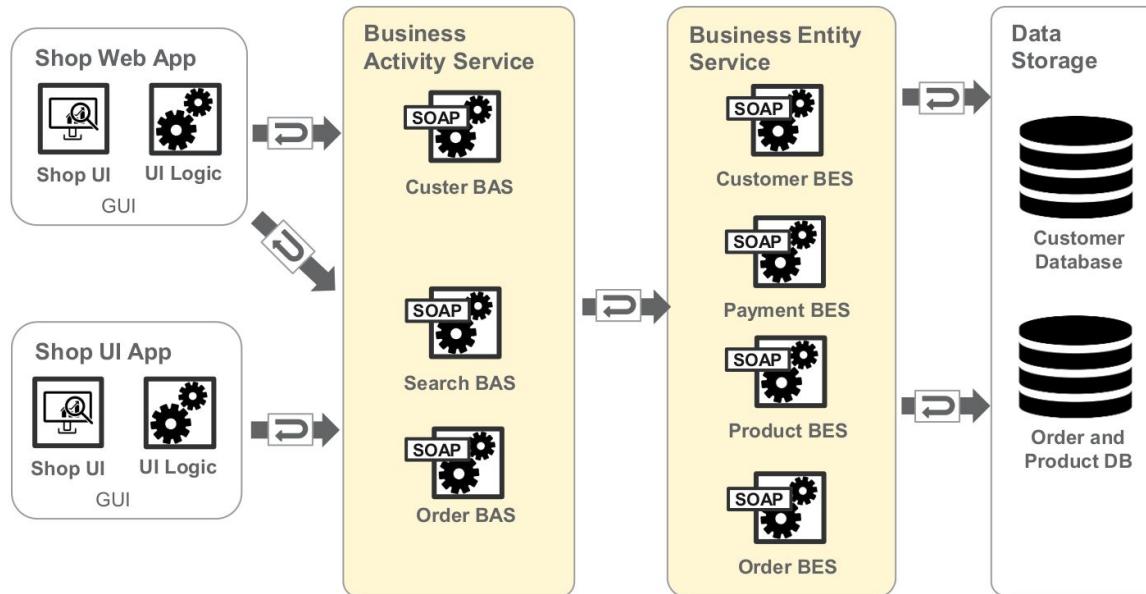
# ■ SOA Approach

Contract-first  
Web Services

Technical layers  
offer their own  
interfaces

Reuse on each  
level

Lower layer  
often wraps  
legacy code

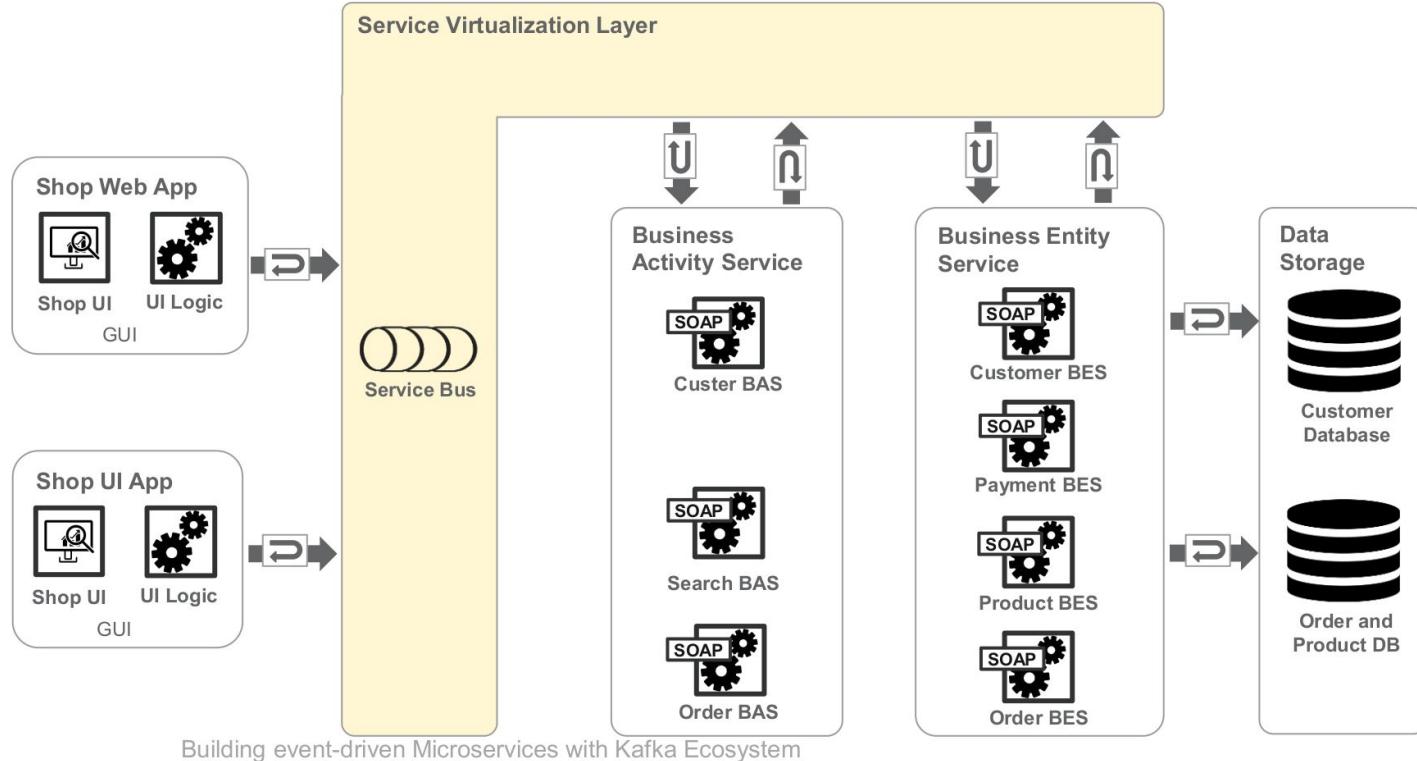


Building event-driven Microservices with Kafka Ecosystem

**trivadis**  
makes IT easier. ■ ■ ■



# ■ Virtualized SOA Approach



# ■ Microservice Approach

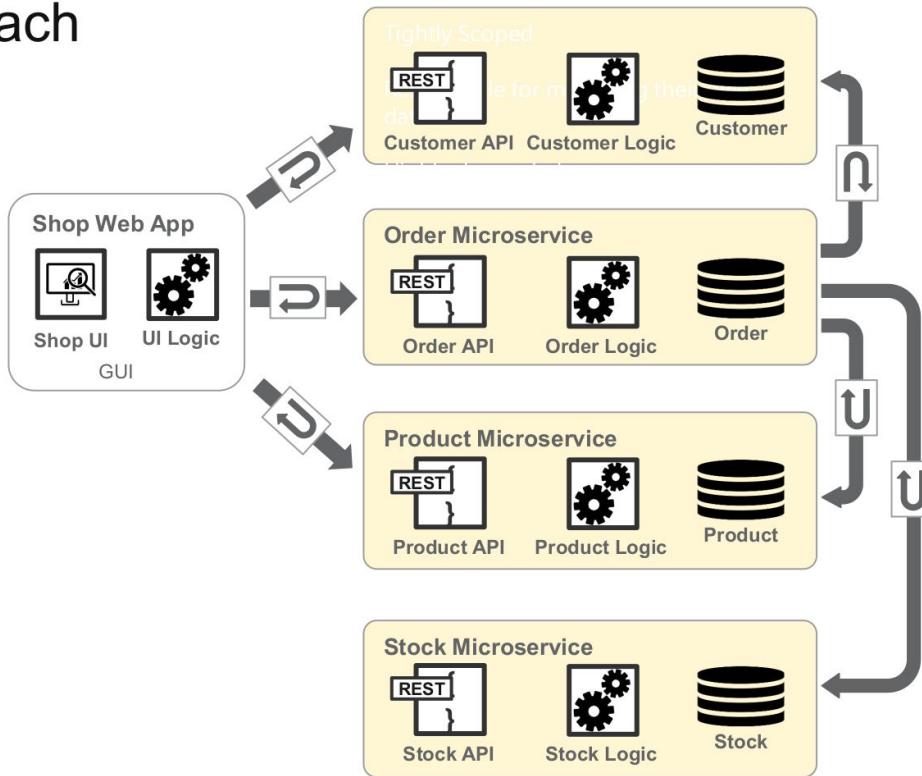
Tightly Scoped behind clear interfaces

Responsible for managing their own data (not necessarily the infrastructure)

Should be highly decoupled

Independently deployable, self-contained and autonomous

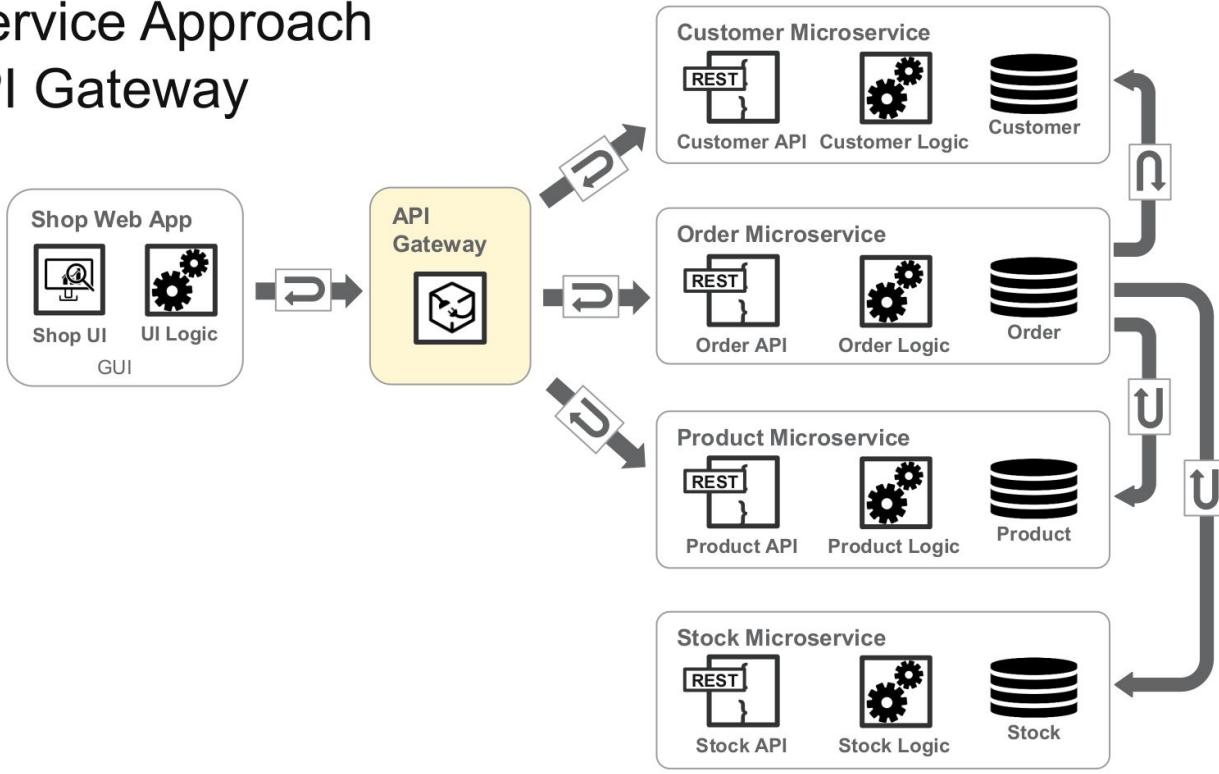
SOA done right ?!



Building event-driven Microservices with Kafka Ecosystem



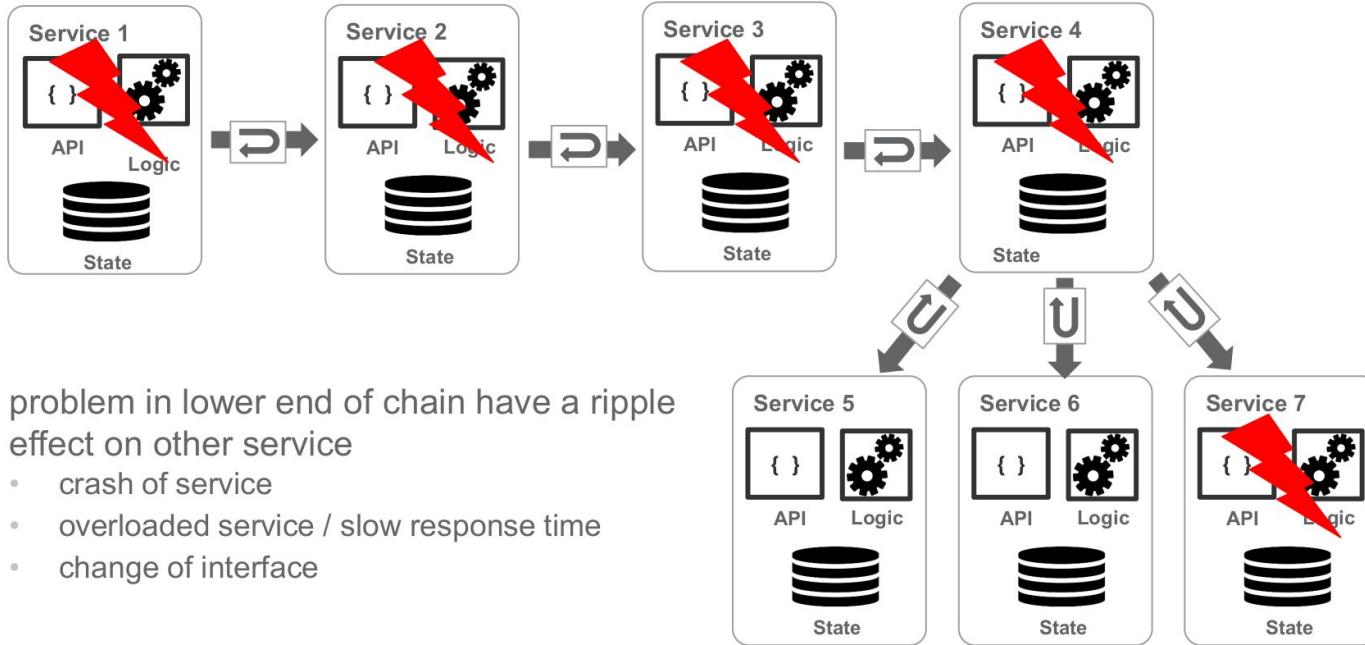
# ■ Microservice Approach with API Gateway



Building event-driven Microservices with Kafka Ecosystem



## ■ Synchronous World of Request-Response leads to tight, point-to-point couplings

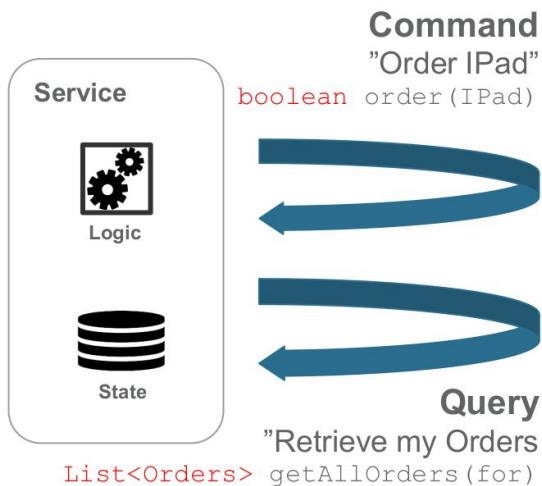


Building event-driven Microservices with Kafka Ecosystem



## ■ Three mechanisms through which services can interact

### Request-Driven (Imperative)



### Event Driven (Functional)

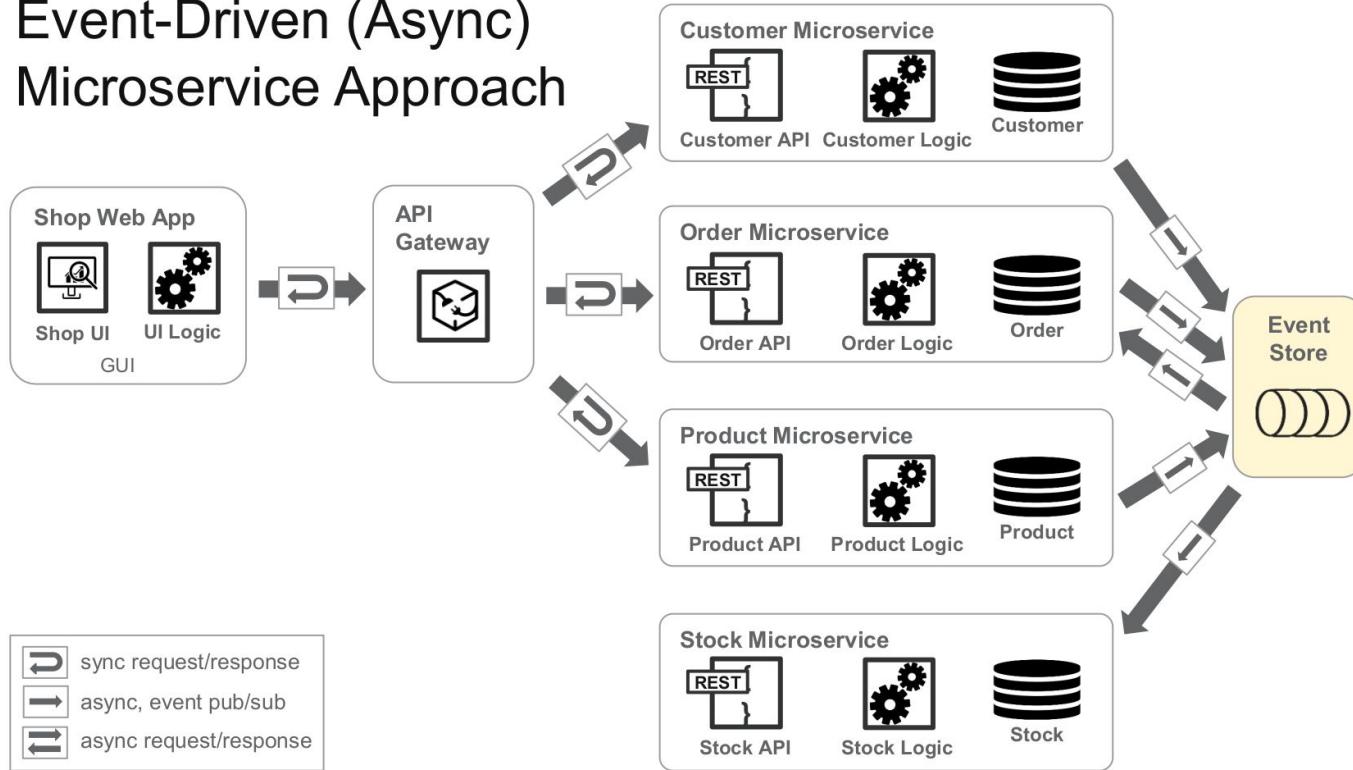


Building event-driven Microservices with Kafka Ecosystem

**trivadis**  
makes IT easier.



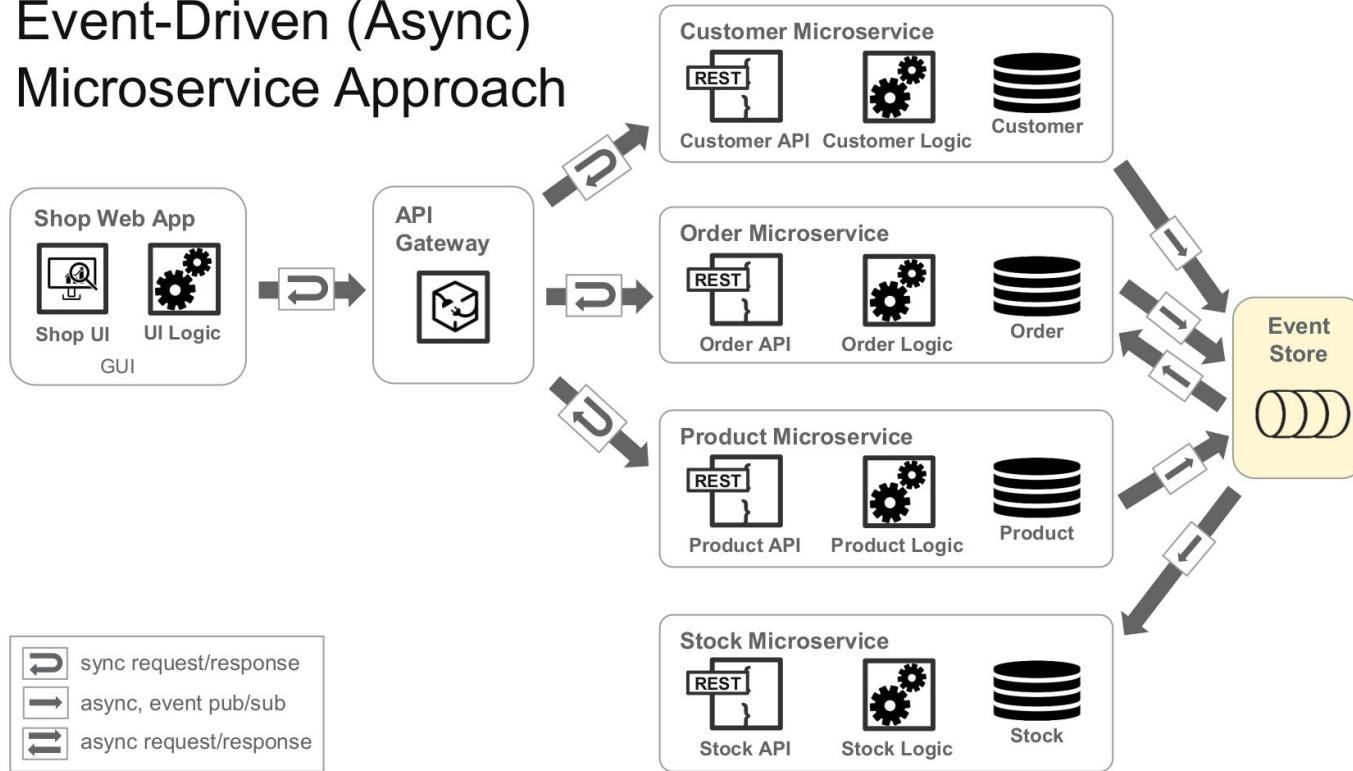
# ■ Event-Driven (Async) Microservice Approach



Building event-driven Microservices with Kafka Ecosystem

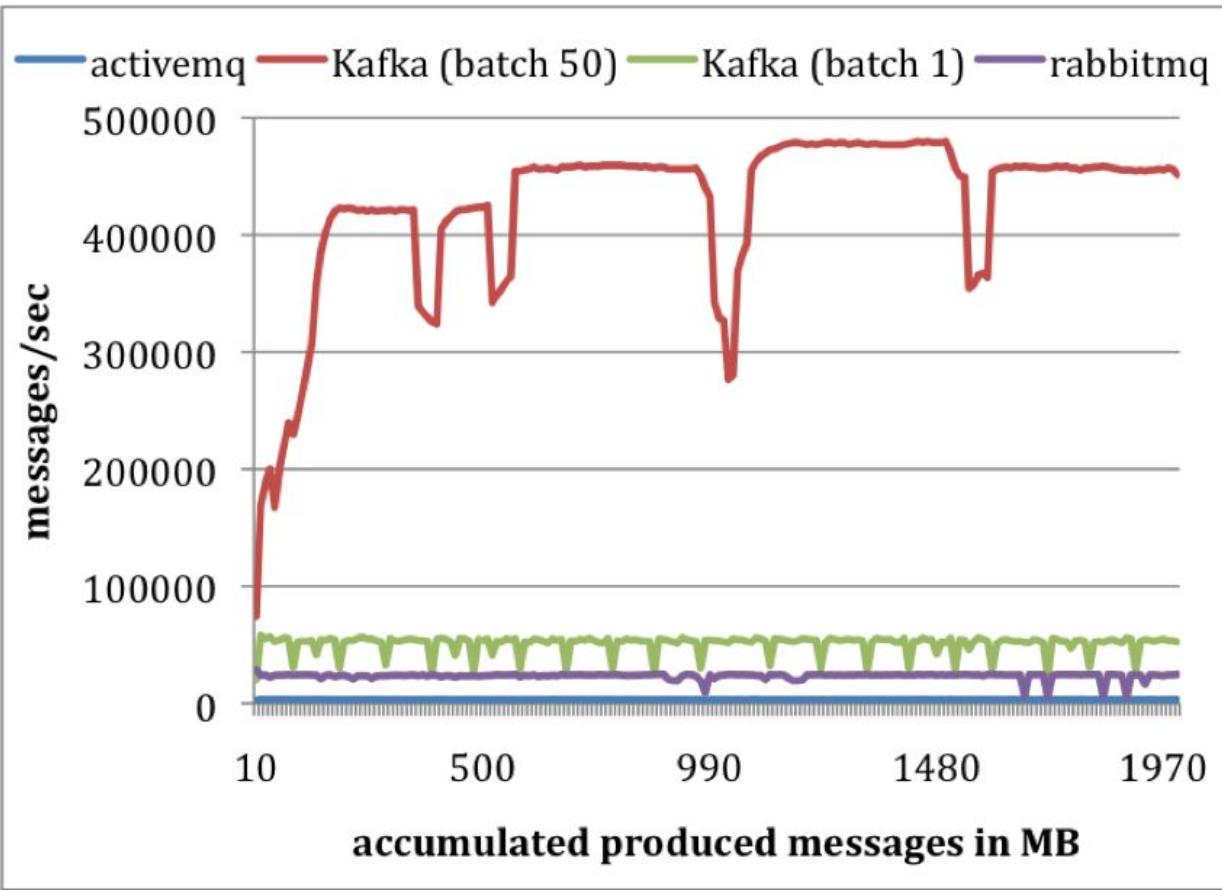


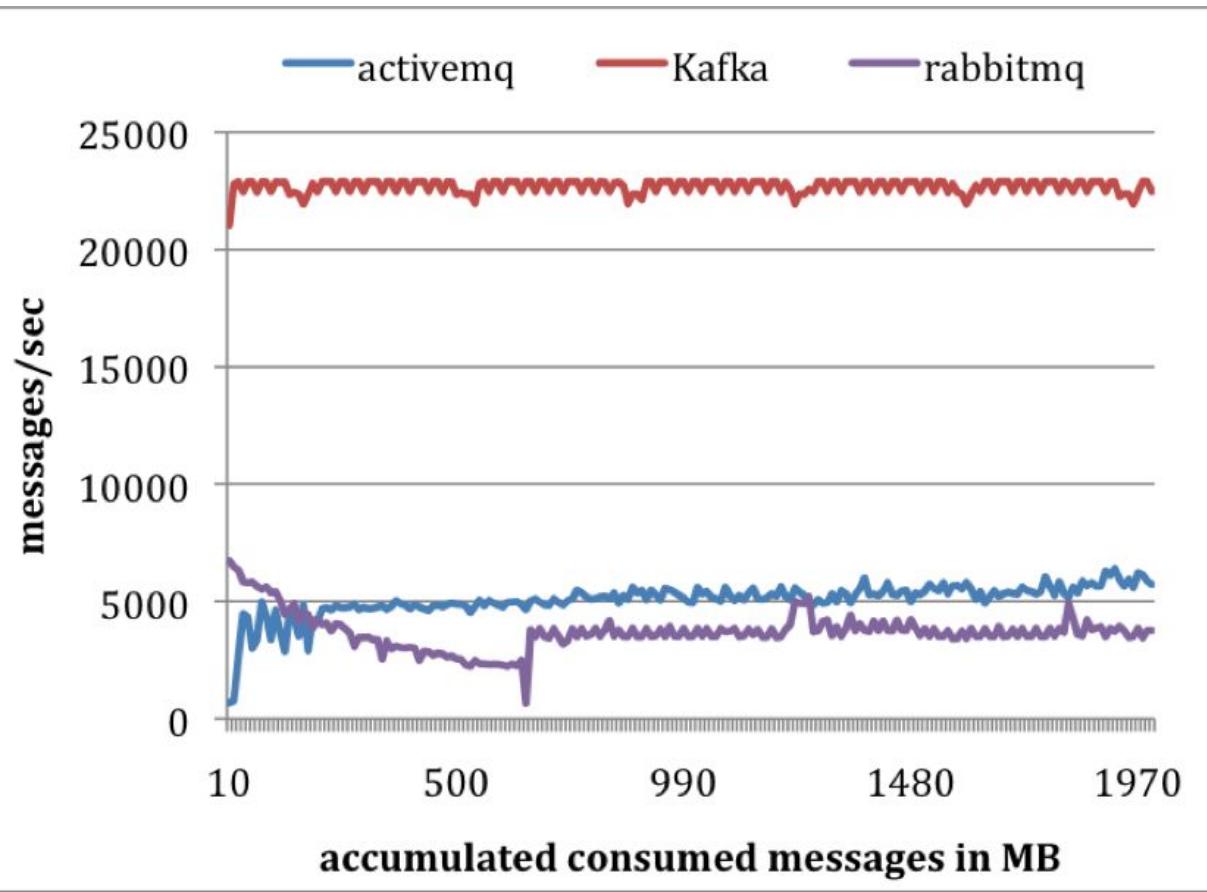
# ■ Event-Driven (Async) Microservice Approach



Building event-driven Microservices with Kafka Ecosystem







# References

- [A Gentle Introduction to Stream Processing](#)
- [Event Sourcing, CQRS, Stream Processing and Apache Kafka: What's the Connection?](#)
- Introduction to Streaming Analytics-Guido Schmutz
- Microservices mit Apache Kafka-Guido Schmutz
- Streaming Visualization-Guido Schmutz

